



Momentum Regime Shifts: A Bayesian Hidden Markov Model Approach to Cross-Sectional Stock Selection

By

Gilad Gang (2148871)

B.Sc. The Hebrew University of Jerusalem

A thesis submitted in partial fulfillment of the requirements for the degree of Master of Science in

Quantitative Finance and Actuarial Science

Tilburg School of Economics and Management
Tilburg University

Supervisor
Denis Kojevnikov

Date: May 2026

Abstract

This thesis investigates how a real-time regime signal interacts with cross-sectional momentum and finds that momentum is not a single stable anomaly but a regime-dependent family of horizon-specific effects: in calm markets, intermediate-horizon continuation dominates the cross-section; in the most acute panic months, beaten-down stocks across all horizons become more likely to outperform, while in milder panic months the cross-section retains continuation structure. A Bayesian Hidden Markov Model produces a filtered probability of market panic, which is combined with momentum at 12 lookback horizons in machine learning models to rank stocks. The regime signal functions as a context variable – it carries no direct linear return-predictive content but conditions how the model jointly uses the full momentum term structure in a given market state. Out-of-sample on 2011–2025 U.S. equities, only a nonlinear model (XGBoost) exploits this structure (Sharpe 1.11, six-factor alpha 24.1%, $t = 4.81$); both a classification-based and a regression-based linear model collapse, establishing that the interaction is inherently nonlinear. Cross-sectional z-score analysis confirms this regime-dependent structure in the raw data. The regime signal is essential (removing it reduces the Sharpe by more than 60%) and most valuable during stress, with the long-short spread widening from 1.3% to 3.1% per month. These findings are consistent with the interpretation that momentum’s well-documented fragility arises partly from applying unconditional models to what is fundamentally a regime-dependent phenomenon.

Acknowledgements

I would like to thank my supervisor, Denis Kojevnikov, for his guidance and feedback throughout this project. I am also grateful to Tilburg University's School of Economics and Management for providing access to WRDS data. Any remaining errors are my own.

Management Summary

Cross-sectional momentum – buying recent winners and selling recent losers – is one of the most profitable and most fragile anomalies in finance. This thesis finds that momentum’s fragility may arise because it is not a single stable pattern but a regime-dependent family of effects: in calm markets, stocks that rose over intermediate horizons continue to outperform; in panic, stocks that have fallen across all lookback horizons become more likely to outperform. The regime signal, derived from a Bayesian Hidden Markov Model estimated on real-time stress indicators, functions as a context variable that conditions how the model jointly uses the full momentum term structure rather than predicting returns directly. Only a nonlinear model (XGBoost) can exploit this structure; linear models with identical features fail entirely, including both a classification-based and a regression-based specification.

The thesis develops a two-stage framework (Section 3). In the first stage, a Bayesian two-state Hidden Markov Model is estimated via Gibbs sampling on four macrofinancial stress indicators (drawdown, dispersion, relative market participation, and credit spread), producing a filtered probability of market panic. In the second stage, this regime signal is combined with momentum at 12 lookback horizons in models of increasing complexity: a deterministic formula, logistic regression, Ridge regression, and XGBoost.

Out-of-sample on 167 months of U.S. equity data (2011–2025), only XGBoost survives long-short construction (Sharpe 1.11, six-factor alpha 24.1%). The regime signal is essential: removing it reduces the Sharpe by more than 60%. Cross-sectional z-score analysis confirms that momentum’s predictive structure shifts with market conditions in the raw data, independent of any model. Stress tests show the strategy can absorb a 12-month recession but is vulnerable to prolonged fundamental deterioration where beaten-down stocks never recover.

Limitations include the single test period (2011–2025), US-only sample, and untested resilience to depression-style downturns. The thesis complements the exposure-scaling approach of Daniel and Moskowitz (2016) with a stock-selection channel: rather than adjusting how much to hold, the regime signal adjusts which stocks to hold.

Contents

1	Introduction	7
2	Literature Review	9
2.1	The Momentum Anomaly and Its Horizon Structure	9
2.2	Momentum Crashes and Exposure-Based Solutions	9
2.3	Regime Models and Regime-Conditional Asset Pricing	10
2.4	Machine Learning in Cross-Sectional Asset Pricing	11
3	Methodology	14
3.1	Stage 1: Regime Classification via Hidden Markov Model	14
3.1.1	Model Specification	14
3.1.2	Bayesian Estimation via Gibbs Sampling	16
3.1.3	Prior Specification	17
3.1.4	Gibbs Sampling Algorithm	19
3.1.5	Implementation and Convergence	23
3.1.6	Regime Identification and Trading Signal	24
3.2	Stage 2: Regime-Conditioned Portfolio Construction	25
3.2.1	Cross-Sectional Stock Selection Methods	26
4	Data	30
4.1	HMM Input Features	30
4.2	Feature Selection	32
4.3	Regime Estimation Results	35
4.4	Cross-Sectional Features	36
5	Results	38
5.1	Portfolio Performance	38
5.1.1	Results Overview	38
5.1.2	Regime Signal Ablation	39
5.1.3	Feature-Set Ablation	40
5.2	How the Model Reorganises Momentum Across Regimes	41
5.2.1	Per-Horizon Momentum SHAP	41
5.2.2	Clustering Selection Modes Across Market States	42
5.2.3	Discussion	49
5.3	Strategy Pitfalls	50
5.3.1	Risk-Aversion Control	51
5.3.2	Stress Test: Prolonged Bear Market	52
5.3.3	Historical Stress: Out-of-Sample Evidence	54

5.4	International Robustness	55
6	Conclusion	57
6.1	Contributions	57
6.2	Limitations	57
6.3	Future Work	59
Appendices		
A	Data Sources and Variable Definitions	66
B	Portfolio Formation, Transaction Costs, and Performance Metrics	67
B.1	Portfolio Construction	67
B.2	Transaction Costs	67
B.3	Performance Metrics	68
C	Feature Importance Methods	69
C.1	SHAP Values for XGBoost	69
D	Alternative Training Targets for XGBoost	70
E	CRRA Dual-Utility Construction	70
F	Robustness Checks	71
F.1	Sub-Period Stability	71
F.2	Turnover and Transaction Cost Sensitivity	72
F.3	Model Sensitivity	73
F.4	Number of Hidden States	74
F.5	January Exclusion	75
G	External Validity	75
G.1	Alternative Train/Test Splits	75
G.2	30-Year Expanding-Window Out-of-Sample Methodology	76
G.3	Factor Model Regressions	76
G.4	Block Bootstrap Inference	77
G.5	Leg-Level CAPM Betas by Regime	78
G.6	Per-cluster Feature Signature ($K = 4$ Partition)	79
G.7	Cluster-Level Cross-Validation	81
G.8	Ridge Regression Baseline	82
G.9	Emission Distribution Robustness	82
H	Label Switching: Sign Correction Procedure	83

I	Convergence Diagnostics	85
I.1	Gelman–Rubin Diagnostic	85
I.2	Seed Convergence of the Regime Signal	86
J	International HMM Feature Selection	87
K	Fundamental Feature Definitions	88

1 Introduction

Cross-sectional momentum, the tendency for recent stock market winners to continue outperforming and losers to continue underperforming, is one of the most robust anomalies in empirical finance (Jegadeesh and Titman, 1993). Yet it is also one of the most fragile: momentum portfolios suffered losses exceeding 50% during the 2009 market rebound (Daniel and Moskowitz, 2016). This tension between long-run profitability and vulnerability to regime shifts motivates the central question of this thesis: how does a real-time regime signal interact with the cross-sectional momentum anomaly, and what does this interaction reveal about momentum’s underlying structure?

Existing responses to momentum’s fragility adjust either exposure or horizon weights based on backward-looking indicators (Section 2). This thesis proposes an alternative: a Bayesian Hidden Markov Model estimated on real-time stress indicators produces a filtered probability of market panic, which enters a machine learning model alongside 12 momentum horizons so that cross-sectional stock selection adapts to the current regime.

The central claim is that this regime signal is not a direct return predictor but a *context variable*: it conditions how the model jointly uses the full momentum term structure in a given market state. No prior study has documented how the full term structure of cross-sectional momentum predictability reorganises across regimes, or that exploiting this reorganisation requires nonlinear modelling.

Out-of-sample evaluation on 167 months of U.S. equity data (2011–2025) yields three main findings. First, only XGBoost produces meaningful long-short returns (Sharpe 1.11, six-factor alpha 24.1%, $t = 4.81$); both linear baselines on identical features collapse (classification-based logistic regression to Sharpe -0.03 and regression-based Ridge to -0.57 ; Appendix G.8), establishing that the interaction is inherently nonlinear rather than an artefact of the training target. Second, the regime signal is essential (removing it reduces the Sharpe by more than 60%) and most valuable during stress, with the long-short spread widening from 1.3% to 3.1% per month in panic. Third, cross-sectional z-score analysis confirms this shift: in calm, the model’s long leg selects stocks with a humped momentum profile peaking at intermediate horizons; in panic, the same leg selects stocks uniformly below average at every horizon from 1 to 12 months.

These findings answer three research questions: (i) does the regime signal’s interaction with momentum require nonlinear modelling, and if so, why does a linear model fail? (ii) how does the cross-sectional structure of momentum change between calm and panic regimes? (iii) how robust is the regime-momentum interaction to alternative signal specifications, risk preferences, and prolonged market stress?

The thesis makes three contributions. First, it documents how the full term structure of cross-sectional momentum predictability reorganises between calm and panic regimes, with short-horizon signals flipping sign during stress. Second, it proposes and validates

a framework in which a macrofinancial regime signal functions as a context variable for stock selection, conditioning how the model jointly weighs the full momentum term structure rather than predicting returns directly. Third, it establishes that this interaction is inherently nonlinear: both classification-based and regression-based linear models fail on identical features, while a tree-based model succeeds.

The remainder of the thesis is organised as follows. Section 2 reviews the momentum anomaly and its horizon structure, momentum crashes and exposure-based risk management, regime models and regime-conditional asset pricing, and machine learning in cross-sectional asset pricing. Section 3 develops the two-stage framework: the Bayesian HMM for regime classification and the three cross-sectional stock selection methods. Section 4 describes the data, HMM feature selection procedure, and cross-sectional features. Section 5 presents the out-of-sample results, including regime-conditional performance, stress tests, SHAP analysis, ablation studies, and the risk-aversion analysis. Section 6 summarises the findings, discusses limitations, and suggests directions for future research.

2 Literature Review

This review covers four strands of literature: the cross-sectional momentum anomaly and its horizon structure, momentum crashes and exposure-based risk management, regime models and their use in regime-conditional asset pricing, and the application of machine learning to cross-sectional return prediction. Each provides a piece of the framework this thesis assembles. None has, on its own, posed the question of how the full term structure of cross-sectional momentum reorganises across regimes.

2.1 The Momentum Anomaly and Its Horizon Structure

Jegadeesh and Titman (1993) established cross-sectional momentum as a pervasive feature of equity markets: portfolios that buy past winners and sell past losers earn significant returns over 3-to-12-month holding periods. The original strategy ranks stocks by cumulative returns over the past twelve months and has been replicated across international markets, asset classes, and sample periods (Jegadeesh and Titman, 2001; Asness et al., 2013).

The economic mechanism behind momentum remains debated. Behavioural explanations include gradual information diffusion (Hong and Stein, 1999) and the joint conservatism-representativeness mechanism of Barberis et al. (1998), while risk-based explanations link momentum to time-varying systematic risk. This thesis does not take a position on the underlying mechanism but exploits the empirical regularity that momentum’s cross-sectional structure varies with market conditions.

A defining feature of momentum is its sensitivity to the lookback horizon. Short lookbacks of one to three months capture short-term reversal effects (Jegadeesh, 1990; Lehmann, 1990), intermediate horizons of six to twelve months capture continuation (Novy-Marx, 2012), and lookbacks beyond twelve months reflect long-term mean reversion (De Bondt and Thaler, 1985). Lee and Swaminathan (2000) document a momentum lifecycle in which the same signal evolves from continuation to reversal as information ages. Together these findings suggest that momentum is a family of horizon-specific patterns whose relative strength varies with market conditions, not a single signal. The framework here tests this directly by exposing all twelve monthly horizons jointly to the cross-sectional model rather than committing to a single lookback ex ante.

2.2 Momentum Crashes and Exposure-Based Solutions

Despite its long-run profitability, momentum is prone to severe crashes. Daniel and Moskowitz (2016) document that momentum portfolios suffered catastrophic losses during the 1932 and 2009 market rebounds, and show that these crashes are predictable: when

the market has recently declined, past losers tend to be high-beta stocks that rally disproportionately during recoveries, while past winners are low-beta defensive stocks that lag. [Barroso and Santa-Clara \(2015\)](#) propose a simple risk-management response: scaling momentum exposure inversely with the strategy’s recent realised volatility, a constant-volatility overlay that substantially raises the Sharpe ratio. [Bianchi et al. \(2022\)](#) pursue a related approach by constructing a crash indicator from the joint time variation of conditional volatility and skewness, again operating on the strategy’s exposure rather than its composition. This thesis explores a complementary composition-side route: rather than scaling exposure during stress, the model reorganises which stocks populate each leg conditional on the regime, with the consequence that the leg-beta inversion underlying the crash mechanism does not occur (Appendix G.5).

The liquidity literature provides a theoretical basis for why momentum crashes reverse. Margin spirals ([Brunnermeier and Pedersen, 2009](#)) and institutional flow pressure ([Vayanos and Woolley, 2013](#)) create temporary price dislocations during stress that revert when selling pressure subsides, the mechanism central to the regime-momentum interpretation developed below.

What unites this strand is the form of remedy: each approach manages momentum’s crash risk by scaling *how much* momentum to hold using backward-looking risk indicators. The framework here addresses the same risk through a different channel, adapting *which stocks* are held using a forward-looking regime signal derived from contemporaneous stress indicators. The two responses are complementary, not substitutes: a stock-selection channel sits alongside the exposure-scaling literature rather than replacing it.

2.3 Regime Models and Regime-Conditional Asset Pricing

Hidden Markov Models have a long history in financial econometrics, originating with [Hamilton \(1989\)](#) for business-cycle analysis and extended to Bayesian estimation via Gibbs sampling by [Albert and Chib \(1993\)](#), with the multivariate-emission generalisation now standard ([Frühwirth-Schnatter, 2006](#); [Hoff, 2009](#)). A common finding in the regime-switching literature for equity markets is that returns are well described by a two-state model with a high-mean, low-volatility state and a low-mean, high-volatility state, where the bear state is less persistent but more volatile ([Ang and Bekaert, 2002](#); [Guidolin and Timmermann, 2007](#)). Hamilton’s original application fitted the regime model to GNP growth rather than returns themselves, a design choice followed here by using macro-financial stress indicators as observable inputs.

A separate strand asks whether asset-pricing anomalies are themselves regime-conditional. [Cooper et al. \(2004\)](#) establish that momentum profits depend on the prior market state, with strong continuation following up-markets and weaker or reversed performance following down-markets. More recent work integrates regime estimation into portfolio con-

struction: [Shu and Mulvey \(2024\)](#) use a sparse jump model to identify factor-specific bull and bear regimes and feed the resulting classifications into a Black-Litterman framework for dynamic factor allocation.

The most directly related work is [Goulding et al. \(2023\)](#), who classify each month into one of four market cycles (Bull, Correction, Bear, Rebound) from the trailing one-month and rolling twelve-month market returns. In each cycle a different blending parameter a determines the weight on one-month versus twelve-month momentum: the composite signal is $a \cdot mom_1 + (1 - a) \cdot mom_{12}$, where $a = 0$ corresponds to pure twelve-month momentum and $a = 1$ to pure one-month momentum. Their dynamic strategy selects the optimal blending parameters for each non-Bull state via grid search on the training sample. The framework here generalises that approach along three axes: from a deterministic four-state classification based on trailing returns to a continuous panic probability from a Bayesian HMM on real-time stress indicators; from two horizons to the full twelve-horizon term structure; and from a parametric linear blend to a learned nonlinear combination. Each of these moves is necessary, individually, to pose a question GHM cannot: whether the cross-sectional structure of momentum reorganises across regimes. Since GHM design their strategies for long-only construction, a direct long-short performance comparison is not appropriate; [Section 5](#) instead uses the GHM market-cycle variable as an alternative regime input within the XGBoost framework, providing a fair test of whether trailing-return-based regime definitions capture the same information as the HMM’s stress indicators.

The unifying feature of this strand is that the regime variable enters portfolio decisions through a hand-coded rule, threshold, or parametric weight, never as a continuous learned feature. This thesis treats π_t^{filter} instead as an input passed to a learner that decides how to use it.

2.4 Machine Learning in Cross-Sectional Asset Pricing

The use of machine learning methods for cross-sectional stock selection has grown rapidly. [Gu et al. \(2020\)](#) compare a wide range of methods on a kitchen-sink characteristic set including multiple momentum horizons, finding that tree-based models and neural networks substantially outperform linear models by capturing nonlinear interactions among firm characteristics. Gradient-boosted trees ([Chen and Guestrin, 2016](#)) have proven particularly effective for tabular financial data because they handle missing values natively, capture nonlinear relationships and interactions without explicit specification, and are robust to irrelevant features. A central challenge in applying these methods to asset pricing is interpretability, addressed here through SHAP values ([Lundberg and Lee, 2017](#); [Lundberg et al., 2020](#)), which decompose each prediction into feature-level contributions and admit an exact polynomial-time algorithm for tree ensembles.

The closest methodological precedent is [Beckmeyer and Wiedemann \(2025\)](#), who use machine learning to learn data-driven weights across past daily returns for a single momentum signal, constructing a dynamically reweighted strategy that outperforms the conventional twelve-minus-one specification. Their finding establishes that data-driven weighting of past returns outperforms fixed schemes. The framework here differs in two respects. First, the weighting is learned not over the daily history of one signal but over a vector of twelve monthly horizons. Second, and more importantly, the weighting is learned *conditional on a regime context variable*, isolating the regime-conditioning role as the primary contribution. The unconditional weighting in [Beckmeyer and Wiedemann \(2025\)](#) averages across calm and panic regimes; this thesis tests whether disaggregating that average reveals a structural reorganisation of the term structure between states.

Despite these advances, no published study examines how the full term structure of cross-sectional momentum predictability reorganises across regimes, or whether exploiting this reorganisation requires nonlinear modelling.

Position of This Thesis

Each strand contributes one ingredient that the others cannot supply on their own. The momentum literature provides the object of study, a family of horizon-specific signals whose relative strength is known to vary with conditions but has been treated unconditionally in standard practice. The crash literature documents the cost of this fragility, shows that exposure-scaling overlays can partially manage it, and motivates a complementary fix at the stock-selection level rather than at the exposure level alone. The regime-conditional asset-pricing literature provides the conditioning variable, although it has consistently been used as a switch or parametric weight rather than as a continuous learned feature. The machine-learning literature provides the function class capable of representing high-dimensional interactions, although its existing applications have addressed kitchen-sink characteristic sets or unconditional momentum weighting rather than regime-conditional momentum structure.

The methodology developed in this thesis combines these four ingredients into a two-stage fix. First, a Bayesian Hidden Markov Model identifies the prevailing regime in real time from macrofinancial stress indicators, producing a continuous filtered panic probability. Second, a tree-based learner consumes this probability alongside the full twelve-horizon momentum vector and learns to select stocks differently in each regime. Both stages are essential: an unreliable regime estimate misinforms the cross-sectional model, and a model that cannot represent regime-conditional shape interactions cannot exploit even a perfect regime estimate. The regime variable is treated as a feature rather than a switch, the momentum signal as a vector rather than a scalar, and the combination as a learned interaction rather than a parametric blend. Each individual move appears some-

where in the prior literature; the targeted combination, applied to the question of how the cross-section reorganises across regimes, does not. Section 3 develops this framework in detail, beginning with the Bayesian HMM and proceeding through the cross-sectional models that consume its output. As anticipated by the regime-as-feature framing, the resulting panic probability is best read as a non-priced state variable, conditioning the joint use of momentum horizons rather than carrying its own risk premium, an interpretation supported by the factor-alpha tests reported in Section 5.

3 Methodology

The empirical framework is two stages. The first stage estimates market regimes using a Bayesian Hidden Markov Model (Section 3.1): a two-state HMM is fit to four macro-financial stress indicators via Gibbs sampling, producing a monthly filtered probability π_t^{filter} that the market is in a panic state. Because this probability conditions only on information available at time t , it is a real-time signal with no look-ahead bias.

The second stage uses this regime signal to guide cross-sectional stock selection (Section 3.2). Three portfolio construction methods of increasing complexity serve distinct analytical roles. The deterministic formula (DET) tests whether simply varying the momentum lookback horizon with market conditions adds value. The logistic regression (LR) tests whether a linear model can exploit the regime signal alongside momentum. XGBoost (XGB) tests whether nonlinear interactions between the regime signal and other features matter. Comparing LR and XGB is the key diagnostic: if the regime signal functions as a context variable rather than a direct return predictor, a linear model should be unable to use it, while a tree-based model should find it important.

3.1 Stage 1: Regime Classification via Hidden Markov Model

3.1.1 Model Specification

This thesis employs a two-state ($K = 2$) Bayesian Hidden Markov Model to classify the market into latent regimes. The choice of $K = 2$ is validated empirically: Appendix F.4 shows that increasing to $K = 3, 4$, or 5 states does not improve downstream portfolio performance and introduces estimation instability. The hidden state $s_t \in \{0, 1\}$ represents the unobserved regime at month t , where $s_t = 0$ denotes the *calm* regime and $s_t = 1$ denotes the *panic* regime. The two states could equally be labelled bull/bear, expansion/recession, or low/high volatility; calm/panic is used throughout this thesis because it reflects the economic nature of the chosen stress features, which are most directly diagnostic of market stress rather than of trend direction or business-cycle phase. Crucially, as emphasized by Hamilton (1989), market regimes are not directly observable: an investor cannot know with certainty whether the market is currently in a calm or panic state, and classifications that look obvious in retrospect are genuinely difficult to identify in real time. The HMM formalises this uncertainty by treating regimes as hidden states that must be inferred probabilistically from observable market data (see Figure 1).

Observation equation Conditional on the regime, the vector of standardized features $\mathbf{z}_t \in \mathbb{R}^D$ (where $D = 4$) is drawn from a regime-specific multivariate Gaussian. The shorthand $\mathbf{z}_{1:t} = (\mathbf{z}_1, \mathbf{z}_2, \dots, \mathbf{z}_t)$ denotes the full history of observations up to and including month t :

$$(\mathbf{z}_t \mid s_t = 0) \sim \mathcal{N}(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0), \quad (\mathbf{z}_t \mid s_t = 1) \sim \mathcal{N}(\boldsymbol{\mu}_1, \boldsymbol{\Sigma}_1). \quad (1)$$

Each regime is characterized by its own mean vector $\boldsymbol{\mu}_k \in \mathbb{R}^D$ and covariance matrix $\boldsymbol{\Sigma}_k \in \mathbb{R}^{D \times D}$, where $k \in \{0 \text{ (calm)}, 1 \text{ (panic)}\}$. The multivariate specification allows the model to capture not only level differences in individual features across regimes, but also changes in the correlation structure; for instance, stress features may become more correlated during panic periods. The Gaussian assumption is standard in the HMM literature. Appendix G.9 verifies its adequacy by re-estimating the model with multivariate Student- t emissions: regime classifications agree on at least 97.8% of months, filtered panic probabilities correlate at ≥ 0.98 , and BIC does not improve, confirming that the Gaussian specification is robust for this application.

Transition equation The evolution of the hidden state follows a first-order Markov chain with a 2×2 transition matrix:

$$\mathbf{P} = \begin{pmatrix} P_{00} & P_{01} \\ P_{10} & P_{11} \end{pmatrix}, \quad (2)$$

where $P_{ij} = P(s_t = j \mid s_{t-1} = i)$ and each row sums to one. The diagonal elements P_{00} (calm persistence) and P_{11} (panic persistence) govern how likely each regime is to continue into the next month. The expected duration of the calm regime is $1/(1 - P_{00})$ months, and likewise $1/(1 - P_{11})$ for the panic regime.



Figure 1: Architecture of the two-state Hidden Markov Model. Hidden states (Calm/Panic) evolve via transition probabilities P_{ij} and emit the four-dimensional observation vector $\mathbf{z}_t$. The forward filter inverts this process to produce $\pi_t^{\text{filter}} = P(s_t = 1 \mid \mathbf{z}_{1:t})$. The lower panel shows an example state sequence illustrating regime persistence and abrupt transitions.

3.1.2 Bayesian Estimation via Gibbs Sampling

The standard frequentist alternative is the Baum–Welch algorithm (Baum et al., 1970), a special case of Expectation-Maximisation (Dempster et al., 1977), which finds maximum likelihood point estimates.

Maximum likelihood estimation of regime-switching HMMs has several known shortcomings for the present application. The likelihood surface is multimodal, with at minimum two equivalent global optima arising from the invariance of the likelihood under state relabelling, and EM can converge to degenerate solutions in which one regime collapses around a single observation (Frühwirth-Schnatter, 2006). Both problems are exacerbated for rare regimes: the panic state in this sample contains only 91 monthly observations in training, where the asymptotic standard errors that frequentist inference relies on are

unreliable and the information matrix can become near-singular at boundary transition probabilities (Hamilton, 1989, footnote 12). The Bayesian approach addresses these issues directly. Posterior sampling integrates over parameter uncertainty rather than relying on a point estimate; conjugate priors regularise rare-regime estimates with explicit, defensible weights (Section 3.1.3); and convergence diagnostics across multiple chains (Section 3.1.5) diagnose multimodality that would otherwise go unnoticed under a single EM run.

The Bayesian approach instead draws from the joint posterior:

$$p(\{s_t\}_{t=1}^T, \{\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k\}_{k=0}^1, \mathbf{P} \mid \mathbf{z}_{1:T}).$$

This posterior is not available in closed form because the hidden states and model parameters depend on each other. Therefore we employ Gibbs sampling (Albert and Chib, 1993; Hoff, 2009), a Markov chain Monte Carlo (MCMC) method, which resolves this by alternating between three conditional updates: (i) the full hidden state path $\{s_t\}_{t=1}^T$, (ii) the emission parameters $(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ for each regime, and (iii) the transition matrix $\mathbf{P}$. Conjugate priors ensure that each block can be sampled exactly from its conditional posterior.

3.1.3 Prior Specification

The priors encode two economic beliefs. First, before seeing the data, neither regime should be forced to have extreme values in any feature; the model should learn the characteristics of calm and panic states from the data. Second, regimes should be persistent rather than switching every few months.

Emission parameters For each regime k , the mean vector and covariance matrix are given a conjugate Normal-Inverse-Wishart (NIW) prior:¹

$$\boldsymbol{\Sigma}_k \sim \mathcal{IW}(\nu_0, \boldsymbol{\Psi}_0), \quad \boldsymbol{\mu}_k \mid \boldsymbol{\Sigma}_k \sim \mathcal{N}\left(\mathbf{m}_0, \frac{\boldsymbol{\Sigma}_k}{\kappa_0}\right). \quad (3)$$

The hyperparameters are chosen as

$$\mathbf{m}_0 = \mathbf{0}, \quad \kappa_0 = 0.01, \quad \nu_0 = D + 2, \quad \boldsymbol{\Psi}_0 = \mathbf{I}_D(\nu_0 - D - 1), \quad (4)$$

where $D = 4$ is the number of features. After observing data, these prior hyperparameters update to posterior counterparts κ_n , $\mathbf{m}_n$, ν_n , and $\boldsymbol{\Psi}_n$ (Equation 9). Each hyperparameter controls a different aspect of the prior belief:

¹The Inverse-Wishart is a probability distribution over positive-definite matrices, making it the natural prior for a covariance matrix: every sample is guaranteed to be symmetric and positive-definite. The Normal-Inverse-Wishart combines this with a conditional Normal prior on the mean. The prior is called *conjugate* because when the data are Gaussian, the posterior takes the same NIW form with updated hyperparameters, making each Gibbs update an exact draw rather than an approximation.

- $\mathbf{m}_0$ is the prior guess for each regime’s mean vector. Setting $\mathbf{m}_0 = \mathbf{0}$ reflects that, since all features are standardized, there is no reason to expect either regime to be centered away from zero before seeing the data.
- κ_0 controls how strongly the prior mean $\mathbf{m}_0$ is trusted relative to the data. The very small value $\kappa_0 = 0.01$ makes this prior nearly uninformative, so the posterior mean is driven overwhelmingly by the observed data.
- ν_0 is the degrees of freedom of the Inverse-Wishart prior, interpretable as the number of pseudo-observations the prior is worth. The value $\nu_0 = D + 2 = 6$ is the smallest integer for which the prior expected value exists, yielding the most diffuse valid prior; a larger choice (e.g., $\nu_0 = 9$) would impose stronger confidence that the covariance is close to $\mathbf{I}_D$. With roughly 150 calm and 91 panic months in training, the posterior degrees of freedom ($\nu_n = \nu_0 + n_k$) far exceed the prior, so the data dominate the covariance estimates.
- Ψ_0 is the scale matrix of the Inverse-Wishart prior. It is calibrated so that the prior expected covariance equals the identity matrix, $\mathbb{E}[\Sigma_k] = \Psi_0/(\nu_0 - D - 1) = \mathbf{I}_D$, implying unit marginal variance and no cross-feature correlation before observing the data.

Transition matrix Each row of the transition matrix is assigned an independent Dirichlet prior:

$$\mathbf{P}_{0,\cdot} \sim \text{Dir}(9, 1), \quad \mathbf{P}_{1,\cdot} \sim \text{Dir}(1, 9). \quad (5)$$

Equivalently, in matrix form,

$$\boldsymbol{\alpha}_{\text{dir}} = \begin{pmatrix} 9 & 1 \\ 1 & 9 \end{pmatrix}.$$

This prior favors persistence in both regimes: absent data, the prior mean transition matrix is

$$\mathbb{E}[\mathbf{P}] = \begin{pmatrix} 0.90 & 0.10 \\ 0.10 & 0.90 \end{pmatrix}. \quad (6)$$

In other words, the prior expects each regime to continue with probability 0.9, corresponding to an expected duration of roughly 10 months. This prior is informative enough to discourage pathological solutions but weak enough to be overridden by the data.

3.1.4 Gibbs Sampling Algorithm

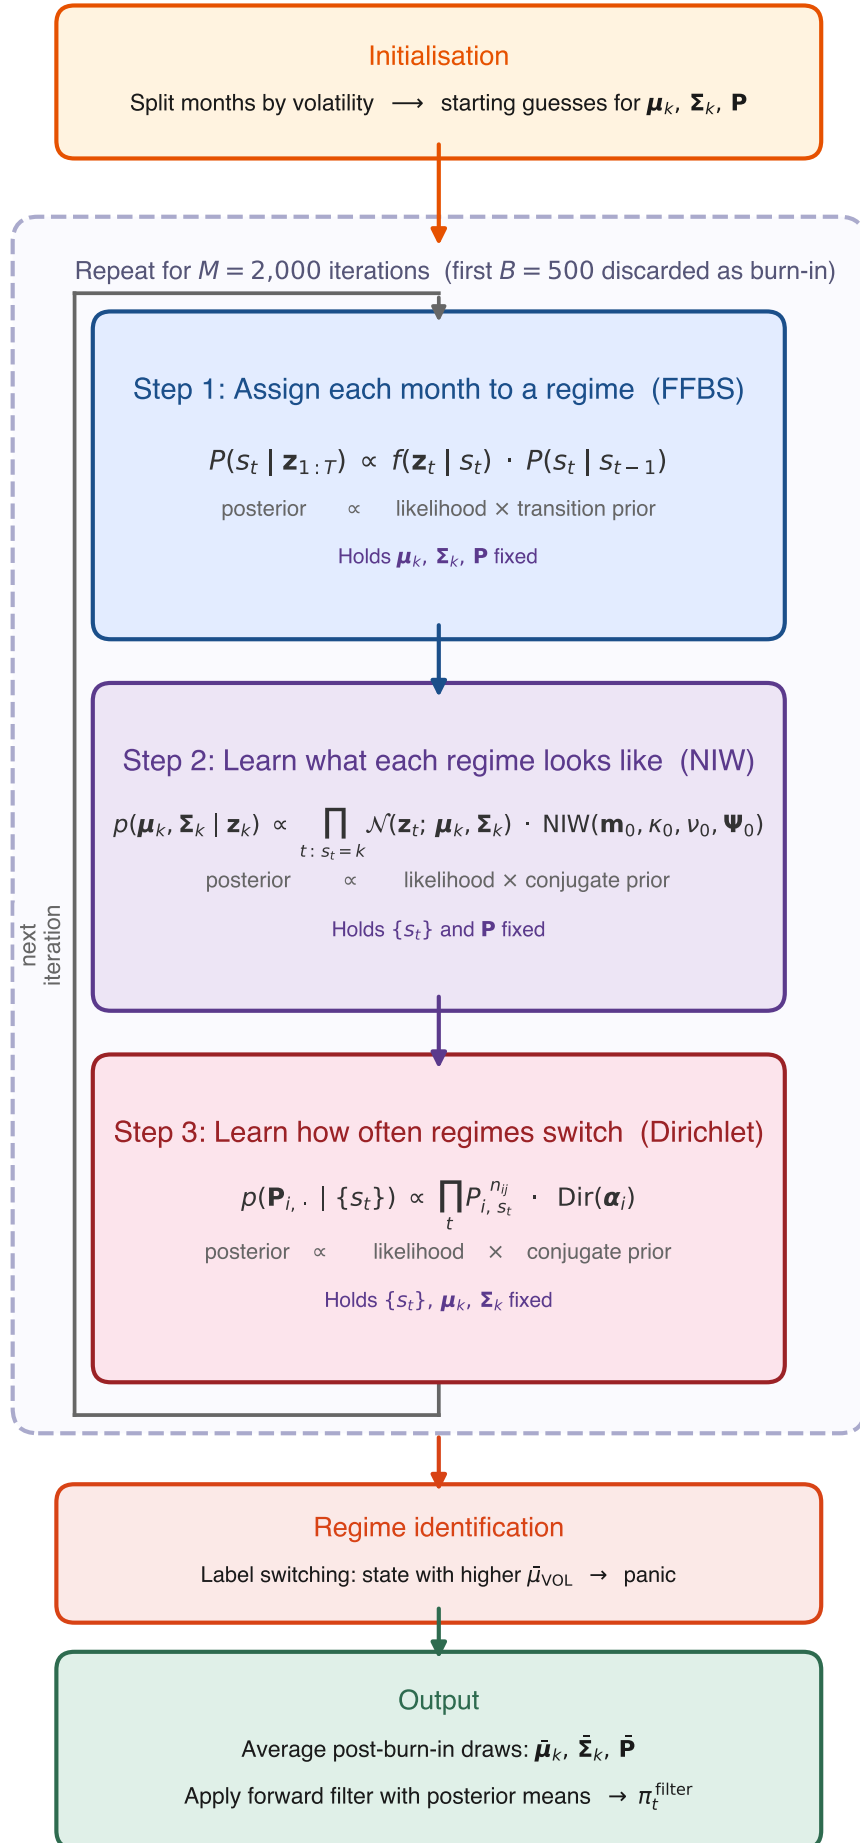
Initialization The Gibbs sampler requires starting values for the emission parameters and the transition matrix. A crude drawdown-based partition is used:² months with below-median drawdown DD_t^z (i.e., deeper drawdowns) are initially assigned to the panic state ($s_t = 1$), and the remaining months to the calm state ($s_t = 0$). The initial emission parameters are then computed from each group:

$$\boldsymbol{\mu}_k^{(0)} = \frac{1}{n_k} \sum_{t: s_t=k} \mathbf{z}_t, \quad \boldsymbol{\Sigma}_k^{(0)} = \text{Cov}(\{\mathbf{z}_t : s_t = k\}),$$

where n_k is the number of months assigned to regime k . The initial transition matrix is set to the prior mean, $\mathbf{P}^{(0)} = \mathbb{E}[\mathbf{P}]$ (Equation 6). These initial parameters give the first forward pass (described below) emission densities that can distinguish the two states.

Iterative updates Each iteration cycles through the three conditional updates described in Section 3.1.2. Figure 2 summarizes the cycle graphically.

²The choice of initialization has negligible impact on the posterior: the sampler is run for 2,000 iterations with 500 discarded as burn-in, and convergence diagnostics (Section 3.1.5) confirm that the chain mixes well regardless of the starting partition.



Step 1: Sampling the hidden state path via FFBS The full latent regime sequence $(s_1, \dots, s_T)$ is sampled using Forward-Filtering Backward-Sampling (FFBS; see Frühwirth-Schnatter, 2006, ch. 11), holding the current emission parameters and transition matrix fixed.

The forward pass computes the filtered probability $\alpha_t(k) = P(s_t = k \mid \mathbf{z}_{1:t})$, the probability of being in state k at time t given all observations so far. This quantity is built in two sub-steps. First, a *prediction step* asks: before seeing this month’s data, how likely is each state? Since the previous state is unknown, the algorithm sums over both possibilities, weighting each by its filtered probability from the previous month and the transition probability:

$$\alpha_{t|t-1}(k) = \alpha_{t-1}(0) P_{0k} + \alpha_{t-1}(1) P_{1k}.$$

Second, an *update step* revises this prediction once the current observation $\mathbf{z}_t$ arrives, using Bayes’ rule: if the observed features look more like a panic month than a calm month, the probability of panic increases. Combining the two steps yields the recursion

$$\alpha_t(k) = \frac{f(\mathbf{z}_t \mid s_t = k) \alpha_{t|t-1}(k)}{f(\mathbf{z}_t \mid s_t = 0) \alpha_{t|t-1}(0) + f(\mathbf{z}_t \mid s_t = 1) \alpha_{t|t-1}(1)}, \quad (7)$$

where $f(\mathbf{z}_t \mid s_t = k) = \mathcal{N}(\mathbf{z}_t; \boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k)$ is the Gaussian emission density.³ The recursion is initialized with $\alpha_0(k) = 1/2$, reflecting no prior preference for either regime.

The forward pass tells us how likely each regime is at every month, but if we simply flipped a weighted coin independently at each t we could produce implausible paths (for example, rapid calm-panic-calm-panic switching even though the transition matrix says regimes persist for roughly ten months). The backward pass prevents this by sampling states in reverse: once a future state has been drawn, the current state is sampled conditional on it, so the entire path respects the Markov transition dynamics.

The backward pass proceeds in two steps:

Terminal state At $t = T$, the filtered probability already conditions on all observations, so the terminal state is drawn directly:

$$s_T \sim \text{Cat}(\alpha_T(0), \alpha_T(1)),$$

³In practice, these densities can be extremely small (or large) for four-dimensional observations, leading to numerical underflow. The implementation avoids this by working in log-space: log-likelihoods $\log f(\mathbf{z}_t \mid s_t = k)$ are computed directly, and the log-sum-exp identity $\log(e^a + e^b) = \max(a, b) + \log(1 + e^{-|a-b|})$ is used whenever probabilities must be normalised.

that is, the algorithm draws $s_T = 0$ (calm) with probability $\alpha_T(0)$ and $s_T = 1$ (panic) with probability $\alpha_T(1)$.⁴

Backward recursion For each earlier time step $t = T - 1, T - 2, \dots, 1$, the algorithm chooses a state consistent with two pieces of information: (i) the filtered belief $\alpha_t(k)$ from the forward pass, and (ii) the already-sampled future state s_{t+1} , which constrains the current state via the transition probability $P_{k,s_{t+1}}$. Combining these two terms and normalising gives:

$$P(s_t = k \mid s_{t+1}, \mathbf{z}_{1:t}) = \frac{\alpha_t(k) P_{k,s_{t+1}}}{\alpha_t(0) P_{0,s_{t+1}} + \alpha_t(1) P_{1,s_{t+1}}}. \quad (8)$$

For example, if the forward filter assigns high probability to calm at time t ($\alpha_t(0)$ is large) and the already-sampled s_{t+1} is also calm (so $P_{0,0}$ is large), then the numerator for $k = 0$ dominates and the algorithm will very likely sample $s_t = 0$.

The output of the full FFBS procedure is one draw from the posterior over state paths. Across Gibbs iterations, different paths are sampled, exploring the full posterior uncertainty.

Importantly, the backward pass is used only during training to improve posterior parameter estimates. The trading signal π_t^{filter} relies solely on the causal forward filter.

Step 2: Updating the emission parameters Given the sampled state path, each regime k has n_k observations assigned to it. From these, two summary statistics are computed: the sample mean $\bar{\mathbf{x}}_k = \frac{1}{n_k} \sum_{t: s_t=k} \mathbf{z}_t$ and the scatter matrix $\mathbf{S}_k = \sum_{t: s_t=k} (\mathbf{z}_t - \bar{\mathbf{x}}_k)(\mathbf{z}_t - \bar{\mathbf{x}}_k)^\top$. The posterior hyperparameters are:

$$\begin{aligned} \kappa_n &= \kappa_0 + n_k, & \mathbf{m}_n &= \frac{\kappa_0 \mathbf{m}_0 + n_k \bar{\mathbf{x}}_k}{\kappa_n}, & \nu_n &= \nu_0 + n_k, \\ \Psi_n &= \Psi_0 + \mathbf{S}_k + \frac{\kappa_0 n_k}{\kappa_n} (\bar{\mathbf{x}}_k - \mathbf{m}_0)(\bar{\mathbf{x}}_k - \mathbf{m}_0)^\top. \end{aligned} \quad (9)$$

Since $\kappa_0 = 0.01 \ll n_k$, the posterior is driven almost entirely by the data.

New regime parameters are then drawn from the updated posterior: $\Sigma_k \sim \mathcal{IW}(\nu_n, \Psi_n)$ and $\boldsymbol{\mu}_k \mid \Sigma_k \sim \mathcal{N}(\mathbf{m}_n, \Sigma_k / \kappa_n)$.

Step 3: Updating the transition matrix The sampled path is scanned to count transitions: $n_{ij} = \sum_{t=2}^T \mathbf{1}(s_{t-1} = i, s_t = j)$ records how many times the path moved from state i to state j . For example, if the sampled path contains 140 calm-to-calm transitions and 10 calm-to-panic transitions, then $n_{00} = 140$ and $n_{01} = 10$. These observed counts

⁴Cat denotes the categorical distribution, the standard notation for a single draw from a discrete probability vector. With two states it reduces to a weighted coin flip.

are added to the Dirichlet prior pseudo-counts to obtain the posterior:

$$\mathbf{P}_{i,\cdot} \mid \{s_t\} \sim \text{Dir}(\alpha_{i0} + n_{i0}, \alpha_{i1} + n_{i1}). \quad (10)$$

For row 0 (calm), the prior pseudo-counts are $(\alpha_{00}, \alpha_{01}) = (9, 1)$ from Equation (5), so the posterior becomes $\text{Dir}(9 + 140, 1 + 10) = \text{Dir}(149, 11)$. The expected persistence probability under this posterior is $149/160 = 0.93$, reflecting that the data reinforced the prior belief in regime persistence. A new transition matrix row is then drawn from this Dirichlet distribution, and the same procedure is applied independently to row 1 (panic).

With all three blocks updated, the sampler returns to Step 1. After sufficient iterations the draws converge to samples from the joint posterior.

3.1.5 Implementation and Convergence

The Gibbs sampler is estimated on the training sample only. To reduce sensitivity to MCMC initialization, the full sampler is run independently with $N_s = 200$ different random seeds. Each chain is run for $M = 2,000$ iterations with the initialization described above. The first $B = 500$ iterations are discarded as burn-in, leaving 1,500 posterior draws per seed.

For each seed, posterior mean parameters $\bar{\mu}_k$, $\bar{\Sigma}_k$, and $\bar{\mathbf{P}}$ are computed by averaging over the post-burn-in draws. A causal forward filter is then applied to the full sample using these parameters to obtain a seed-specific filtered probability $\pi_t^{\text{filter},(s)}$. The final trading signal is the average across seeds:

$$\pi_t^{\text{filter}} = \frac{1}{N_s} \sum_{s=1}^{N_s} \pi_t^{\text{filter},(s)}. \quad (11)$$

This Bayesian model averaging is a methodological requirement, not a performance optimisation: with a single Gibbs chain the out-of-sample Sharpe has mean 1.04 and interquartile range 0.10 across random HMM seed draws, while averaging 50 chains raises the mean to 1.09 and collapses the IQR to 0.02 (Appendix I.2). The 1.11 figure reported in Section 5 therefore reflects the production 200-seed configuration and should be interpreted as a posterior-averaged quantity rather than a single-chain point estimate.

Convergence Diagnostics Convergence is assessed in two ways. First, trace plots of the regime means and persistence probabilities are visually inspected to verify that the chain mixes around a stable posterior region after burn-in. Second, effective sample sizes (ESS) are computed for key parameters. Because successive MCMC draws are autocorrelated, n draws from the chain contain less information than n independent samples. The ESS quantifies the number of truly independent draws the chain is equivalent to. For a

post-burn-in chain of length n ,

$$ESS = \frac{n}{1 + 2 \sum_{\ell=1}^{L^*} \hat{\rho}(\ell)}, \quad (12)$$

where $\hat{\rho}(\ell)$ is the estimated autocorrelation at lag ℓ ,⁵ and L^* is the first lag at which the autocorrelation becomes negative. High ESS indicates that the chain provides many effectively independent posterior draws.⁶

While ESS measures efficiency within a single chain, it cannot detect whether different chains have converged to different modes. The Gelman–Rubin $\hat{R}$ statistic (Gelman et al., 2013) addresses this by comparing within-chain and between-chain variance across M independent chains, each of length n . Let s_m^2 denote the variance of chain m and $\bar{\theta}_m$ its mean. The within-chain variance is:

$$W = \frac{1}{M} \sum_{m=1}^M s_m^2, \quad (13)$$

and the between-chain variance is:

$$B = \frac{n}{M-1} \sum_{m=1}^M (\bar{\theta}_m - \bar{\theta})^2, \quad (14)$$

where $\bar{\theta} = \frac{1}{M} \sum_m \bar{\theta}_m$. If all chains have converged to the same posterior, B should be small relative to W . The diagnostic is:

$$\hat{R} = \sqrt{\frac{\hat{V}}{W}}, \quad \hat{V} = \frac{n-1}{n} W + \frac{B}{n}. \quad (15)$$

Values of $\hat{R}$ near 1.0 indicate convergence; the standard threshold is $\hat{R} < 1.1$. To evaluate convergence, five of the 200 chains are selected at random. For the selected feature set (DD+DISP+REL_N+CS), all parameters have $\hat{R} < 1.01$ (Appendix I.1).

Trace plots of the regime means and persistence probabilities (Figure 12, Appendix I) provide visual confirmation that all chains converge to stable posterior values after the burn-in period.

3.1.6 Regime Identification and Trading Signal

Label Switching and Regime Naming In a two-state HMM, the labels 0 and 1 are intrinsically arbitrary: relabeling the states leaves the likelihood unchanged. To identify

⁵ $\hat{\rho}(\ell) = \frac{1}{(n-\ell)\hat{\sigma}^2} \sum_{i=1}^{n-\ell} (\theta_i - \bar{\theta})(\theta_{i+\ell} - \bar{\theta})$, where θ_i is the i -th post-burn-in draw, $\bar{\theta}$ its sample mean, and $\hat{\sigma}^2$ its sample variance.

⁶An ESS of at least 100 per parameter is a common minimum for reliable posterior inference (Gelman et al., 2013).

which state corresponds to panic, the known crisis-period behavior of each feature is used.⁷ Two features rise during crises (DISP: +1, CS: +1) and two fall (DD: -1, REL_N: -1). The regime whose posterior mean best aligns with these crisis directions is designated as “panic”; if the Gibbs sampler assigns the panic characteristics to state 0 rather than state 1, the labels are swapped so that π_t^{filter} always represents the probability of the panic state.

Two caveats follow from this naming convention. The HMM identifies two regimes by clustering on the emission features; the label “panic” is then assigned to the cluster whose posterior mean aligns with crisis directions. The label therefore reflects a statistical resemblance to historical crisis periods rather than a directly observed market state. Relatedly, the economic interpretations used later in the thesis (liquidity pressure, constrained arbitrage) are readings of what this cluster *likely* corresponds to given the features it loads on, not direct identification claims about the underlying mechanism.

Filtered Regime Probability Using the posterior mean parameters $\bar{\mu}_k$, $\bar{\Sigma}_k$, and $\bar{\mathbf{P}}$ from Section 3.1.5, the forward recursion (Equation 7) is applied to the full sample to obtain:

$$\pi_t^{\text{filter}} = P(s_t = \text{panic} \mid \mathbf{z}_{1:t}; \bar{\mu}, \bar{\Sigma}, \bar{\mathbf{P}}). \quad (16)$$

A value near 0 indicates high confidence in calm, near 1 indicates panic, and intermediate values reflect genuine ambiguity. This filtered probability serves as the regime signal passed to the portfolio construction methods in Section 3.2.1. Figure 4 and Table 4 in Section 4 visualise the resulting signal and confirm that the two regimes are meaningfully distinct.

3.2 Stage 2: Regime-Conditioned Portfolio Construction

The filtered regime probability π_t^{filter} from Stage 1 provides a monthly measure of market stress, but on its own it does not select stocks. This stage translates the regime signal into cross-sectional portfolios. Three cross-sectional stock selection methods of increasing complexity are developed: a deterministic regime-adaptive momentum formula (DET), a logistic regression classifier (LR), and an XGBoost gradient-boosted tree regressor (XGB). Each method combines π_t^{filter} with momentum signals at multiple horizons to produce a monthly score for every stock, which is then used to rank stocks into a long-short portfolio: the top decile (using NYSE breakpoints) forms the long leg, the bottom decile forms the short leg, and both legs are value-weighted. The stock-level inputs (momentum horizons and the regime signal) are described in Section 4.

A key distinction from standard momentum strategies is worth emphasising. In the

⁷The sign directions are determined by a data-driven procedure comparing crisis and non-crisis 95th percentiles; see Appendix H for details.

traditional approach (Jegadeesh and Titman, 1993), stocks are ranked directly by their raw trailing return at a fixed lookback, so the long leg always holds the stocks with the *highest* momentum and the short leg holds the *lowest*. Here, by contrast, stocks are ranked by their predicted forward return, a learned function of multiple momentum horizons and the regime signal. Because the model learns which momentum *level* is associated with future outperformance in a given market state, the long leg need not contain high-momentum stocks at all: in deep crisis, the model may assign the highest predicted returns to stocks whose momentum is below the cross-sectional average at every horizon (Section 5.2.2). The portfolio construction remains identical (top and bottom deciles), but the criterion for entering each decile is regime-dependent rather than fixed.

The deterministic formula serves as a transparent baseline that tests whether simply varying the momentum lookback horizon with market conditions adds value. The two machine learning methods go further by jointly learning how momentum at multiple horizons and the regime signal interact to predict forward returns. We evaluate every method in long-short construction rather than long-only, for three reasons. First, a long-short portfolio nets out market beta on the two sides, so the remaining return reflects the cross-sectional ranking itself rather than a levered market bet. Second, the regime mechanism this thesis studies is two-sided: in panic the model predicts reversal, meaning it identifies stocks expected to fall as well as stocks expected to rise, and the short leg directly tests whether those downward predictions are informative. Long-only construction would discard that test. Third, long-short is the standard construction in the momentum literature (Jegadeesh and Titman, 1993; Daniel and Moskowitz, 2016; Barroso and Santa-Clara, 2015), which makes our Sharpe ratios and factor alphas directly comparable to existing results. The results that follow (Section 5) not only compare the out-of-sample performance of the three portfolios but also use SHAP values and coefficient analysis to examine *which* features drive the predictions and *how* their importance shifts across calm and panic regimes. This feature-level analysis provides economic insight into the signals that distinguish regime-conditioned stock selection from standard momentum.

3.2.1 Cross-Sectional Stock Selection Methods

DET: Deterministic Regime-Adaptive Formula The simplest approach exploits the economic intuition that the optimal momentum lookback horizon varies with market conditions. The lookback horizon and stock score are:

$$\ell_t = \text{clip}(\text{round}(12 - 11 \cdot \pi_t^{\text{filter}}), 1, 12), \quad (17)$$

$$\text{score}_t^s = \prod_{j=1}^{\ell_t} (1 + r_{t-j}^s) - 1. \quad (18)$$

Stocks are then ranked by this score; the top decile forms the long leg and the bottom decile the short leg, as described in Appendix B.1.

When $\pi_t^{\text{filter}} \approx 0$ (calm), $\ell_t = 12$, recovering the standard 12-month momentum signal that captures continuation in trending markets (Jegadeesh and Titman, 1993). When $\pi_t^{\text{filter}} \approx 1$ (panic), $\ell_t = 1$, switching to 1-month reversal to exploit the tendency of oversold stocks to recover after panic selling (Goulding et al., 2023). At intermediate values, the lookback smoothly interpolates between these extremes.

LR: Logistic Regression A logistic regression classifier is trained to predict whether a stock’s forward return exceeds the cross-sectional median in each month. Classification is chosen over direct return regression because monthly stock returns are extremely noisy, and predicting the binary outcome “above or below median” is a more stable learning target that still produces a useful ranking of stocks. Classification targets have been explored as an alternative to regression in cross-sectional asset pricing, though regression targets generally produce better portfolio performance (Gu et al., 2020); the Ridge regression baseline (Section 3.2.1) confirms that the choice of target does not drive the linear model’s failure.⁸

$$P(y_t^s = 1 \mid \mathbf{x}_t^s) = \sigma(\boldsymbol{\beta}^\top \tilde{\mathbf{x}}_t^s), \quad (19)$$

where $y_t^s = \mathbb{1}(r_{t+1}^s > \text{median}_s\{r_{t+1}^s\})$ is a binary label indicating whether stock s outperforms the cross-sectional median in the following month. The feature vector $\tilde{\mathbf{x}}_t^s = (mom_{1,t}^s, \dots, mom_{12,t}^s, \pi_t^{\text{filter}})$ contains the 12 momentum lookbacks and the regime signal, standardized using training-set means and standard deviations. The model is estimated by maximizing the ℓ_2 -regularized log-likelihood:

$$\hat{\boldsymbol{\beta}} = \arg \max_{\boldsymbol{\beta}} \left[\sum_i \left(y_i \log \sigma(\boldsymbol{\beta}^\top \tilde{\mathbf{x}}_i) + (1 - y_i) \log(1 - \sigma(\boldsymbol{\beta}^\top \tilde{\mathbf{x}}_i)) \right) - \frac{1}{2C} \|\boldsymbol{\beta}\|_2^2 \right], \quad (20)$$

with regularization strength $C = 1.0$. The ℓ_2 penalty shrinks coefficients toward zero to guard against overfitting given the large number of correlated momentum features. Once $\hat{\boldsymbol{\beta}}$ is estimated on the training sample, the stock score for any observation is obtained by plugging $\hat{\boldsymbol{\beta}}$ back into the prediction equation:

$$\text{score}_t^s = \sigma(\hat{\boldsymbol{\beta}}^\top \tilde{\mathbf{x}}_t^s) = P(y_t^s = 1 \mid \mathbf{x}_t^s; \hat{\boldsymbol{\beta}}). \quad (21)$$

Stocks with higher predicted probability of beating the cross-sectional median receive higher scores and are more likely to enter the long leg; those with the lowest scores enter the short leg.

⁸The logistic sigmoid function $\sigma(z) = 1/(1 + e^{-z})$ maps the linear combination $\boldsymbol{\beta}^\top \tilde{\mathbf{x}}$ from the real line to the interval $(0, 1)$, so the output can be interpreted as a probability.

XGB: XGBoost Regressor An XGBoost gradient-boosted tree regressor (Chen and Guestrin, 2016) is trained to predict the forward monthly return directly. Unlike LR, which classifies stocks as above or below the median, XGBoost predicts the forward return directly. This preserves magnitude information: a stock expected to gain 10% ranks higher than one expected to gain 2%, a distinction that binary classification discards. The feature vector for each stock-month observation comprises the 13 features described in Section 4.4: the 12 momentum lookbacks ($mom_1, \dots, mom_{12}$) and the regime signal (π_t^{filter}). By including all horizons jointly, XGBoost can learn how the full momentum term structure interacts with the regime signal, for example by conditioning its joint use of all 12 horizons on the level of π_t^{filter} in ways that cannot be decomposed into per-horizon slope changes.

XGBoost builds an additive ensemble of M regression trees, where each successive tree fits the residuals of the current ensemble:

$$F(\mathbf{x}) = \sum_{m=1}^M \eta \cdot h_m(\mathbf{x}), \quad (22)$$

where h_m is the m -th decision tree and η is the learning rate that shrinks each tree’s contribution to prevent overfitting. Each tree h_m is grown to minimize a regularized objective combining the squared-error loss on the current residuals with penalty terms on tree complexity (number of leaves and leaf weights). The predicted return $\hat{r}_{t+1}^s = F(\mathbf{x}_t^s)$ serves as the stock score. The model is trained with $M = 500$ trees, a learning rate of $\eta = 0.05$, maximum tree depth of 4, and row and column subsampling rates of 0.8. To reduce sensitivity to random initialization, an ensemble of 50 independent XGBoost models (each with a different random seed) is trained; their predictions are averaged before ranking stocks into portfolio deciles.

A maximum depth of 4 allows the model to capture up to four-way feature interactions. The choice of depth is consequential: depth 1 permits no interactions between features, depth 2 permits only pairwise interactions (e.g., one split on π_t^{filter} followed by one split on a single momentum horizon), and depth ≥ 3 permits higher-order conditioning across multiple horizons simultaneously. Table 19 reports the depth sweep: annualised return rises monotonically from depth 1 (8.7%, Sharpe 0.53) through depth 4 (21.7%, Sharpe 1.11), with the largest gains at depths 2 and 3 where the model first gains multi-way conditioning. A depth-2 model, which conditions on π_t^{filter} and one momentum horizon per tree, achieves 15.7%, 28% below the full model. Volatility is nearly constant across depths ($\approx 19.5\%$): deeper trees improve signal quality rather than amplifying risk-taking. The $K = 4$ cluster decomposition in Section 5.2.2 provides a specific account of why depth 4 is exactly the right budget: distinguishing the four multi-horizon shapes within a π_t^{filter} -conditional tree branch requires depth 1 for the regime gate plus at least depth 3 for within-regime shape reading. Performance peaks at depth 4 and declines beyond depth 5.

The value of the regime signal is architecture-dependent: a depth-matched Random Forest on the same 13 features achieves only Sharpe 0.73 and demotes π_t^{filter} to 18% of total SHAP (versus 45% for XGBoost; Appendix F.3). XGBoost is chosen because its sequential residual fitting concentrates splits on whichever feature most reduces error, turning the regime signal into the model’s primary gating variable rather than one feature among thirteen.

Isolating Linearity from the Training Target LR and XGB differ not only in functional form (linear vs. nonlinear) but also in their training targets: LR predicts a binary classification label, while XGB predicts the continuous forward return. To ensure that any performance difference reflects linearity rather than the choice of target, a Ridge regression baseline is included. Ridge regression is a linear model (identical functional form to logistic regression) that predicts continuous forward returns (identical target to XGBoost), using the same 13 standardised features:

$$\hat{r}_{t+1}^s = \boldsymbol{\beta}^\top \tilde{\mathbf{x}}_t^s, \quad \hat{\boldsymbol{\beta}} = \arg \min_{\boldsymbol{\beta}} \left[\sum_i (r_{i,t+1} - \boldsymbol{\beta}^\top \tilde{\mathbf{x}}_i)^2 + \alpha \|\boldsymbol{\beta}\|_2^2 \right], \quad (23)$$

with $\alpha = 1.0$. Results are robust across regularisation strengths from $\alpha = 0.01$ to $\alpha = 1,000$ (Appendix G.8). If Ridge regression succeeds where logistic regression fails, the classification target is the bottleneck; if both linear models fail, linearity itself is the constraint.

The Role of π_t^{filter} A key distinction between the linear models (LR and Ridge) and XGB lies in how each model uses the regime signal. A linear model assigns π_t^{filter} a single coefficient, which is near zero because the regime signal is not a direct return predictor. XGBoost, by contrast, uses π_t^{filter} as a context variable through tree splits, conditioning its joint use of all momentum horizons on the current market state. This conditioning operates through multi-dimensional feature interactions (Gu et al., 2020) and cannot be represented by a single linear coefficient or decomposed into per-feature slope changes. The empirical consequences of this distinction are analysed in Section 5.2.

With the two-stage framework in place, the next section describes the data, sample construction, and the feature selection procedure used to choose the HMM’s input features.

4 Data

This section describes the data and features that implement the two-stage framework of Section 3. The HMM features are chosen so that the resulting regime signal can serve as a context variable for momentum’s term structure; the cross-sectional features provide the full set of momentum horizons whose regime-dependent predictability the models must learn. The stock-level sample spans January 1990 to November 2025 and includes all ordinary common shares (CRSP share codes 10 and 11) listed on the NYSE, AMEX, or NASDAQ with a price above \$1. The \$1 price floor admits small and illiquid stocks into the universe, but NYSE breakpoints for decile classification tilt portfolio composition toward stocks large enough to be actively traded, and value-weighting within each leg further reduces the influence of the smallest names. Following [Shumway \(2001\)](#), delisting returns are incorporated to prevent survivorship bias: actual delisting returns are used when available, and -30% is imputed for performance-related delistings with missing returns.⁹ The sample is split into a training period (1990–2010) used to estimate all models and a test period (2011–2025, 167 months) used exclusively for out-of-sample evaluation; no model parameters are re-estimated during the test period. Table 1 provides an overview.

	Train (1990–2010)	Test (2011–2025)	Full
Unique stocks	13,363	7,037	16,657
Stock-month observations	1,161,014	557,087	1,718,101
Avg. stocks per month	4,838	3,336	4,221
Market-level months	240	167	407

Table 1: Sample summary. The sample includes all ordinary common shares listed on NYSE, AMEX, or NASDAQ with price above \$1.

4.1 HMM Input Features

The HMM classifies market regimes using four monthly macrofinancial stress indicators. Each captures a distinct dimension of market stress:

1. **DD (Drawdown):** The percentage decline of the cumulative market price index from its rolling 12-month peak:

$$DD_t = \frac{P_t - \max(P_{t-11}, \dots, P_t)}{\max(P_{t-11}, \dots, P_t)}. \quad (24)$$

DD directly measures the depth of ongoing market declines. [Daniel and Moskowitz \(2016\)](#) show that momentum crashes are predictable from past market declines, making DD the natural anchor for regime identification.

⁹Full data sources and detailed variable definitions are provided in Appendix A.

2. **DISP (Return Dispersion):** The log cross-sectional standard deviation of individual stock returns:

$$DISP_t = \log(\text{std}_s(r_{s,t})). \quad (25)$$

DISP captures cross-sectional *disagreement* among stocks. During stress, individual stocks move in different directions rather than together, causing dispersion to spike.

3. **REL_N (Relative Market Participation):** The log ratio of the number of actively traded stocks to its trailing 12-month average:

$$REL_N_t = \log\left(\frac{N_t}{\frac{1}{12} \sum_{\ell=1}^{12} N_{t-\ell}}\right). \quad (26)$$

REL_N turns negative during crises when stocks delist, halt trading, or are acquired at elevated rates. It captures structural market damage that the other features do not measure.

4. **CS (Credit Spread):** The difference between Moody's BAA and AAA corporate bond yields:

$$CS_t = Y_t^{\text{BAA}} - Y_t^{\text{AAA}}. \quad (27)$$

The credit spread provides an independent signal from the corporate bond market, helping the HMM distinguish genuine stress episodes from equity-specific turbulence.

All four features are winsorised at the 1st and 99th percentiles (i.e., values below the 1st or above the 99th percentile are capped at those bounds to limit the influence of outliers) using training-sample statistics, and standardised to zero mean and unit variance.




Figure 3: Time series of the four HMM input features (DD, DISP, REL_N, CS), standardised using training-sample statistics. Shaded regions indicate NBER recession periods.

4.2 Feature Selection

The regime signal’s ability to capture the shift in momentum’s cross-sectional structure depends critically on which stress indicators the HMM observes. The selection problem has two constraints. First, market drawdown (DD) must be included: it is the most economically direct measure of the mechanism identified by [Daniel and Moskowitz \(2016\)](#) (momentum crashes follow market declines) and consistently emerges as the dominant regime-identifying feature. Second, the number of features is capped at four. With 241 training months split roughly 150/91 between calm and panic regimes, each regime has sufficient observations to estimate a 4-dimensional mean vector and covariance matrix (10 free covariance parameters per state). Increasing beyond four features risks estimation instability.

Candidate pool Nine monthly macrofinancial indicators, all available from 1990 onward, are considered:

- *Equity market*: market drawdown (DD), realised volatility (VOL), cross-sectional return dispersion (DISP), relative market participation (REL_N), advance-decline ratio (ADR), and cross-sectional return skewness (SKEW).

- *Credit and macro*: BAA–AAA credit spread (CS), log VIX (LVIX), and the 10-year minus 2-year Treasury term spread (TERM).

All candidates are observable in real time at the end of each month and standardised to z -scores using training-period statistics.

Selection procedure The selection proceeds through three passes of increasing rigour. In the first pass (*quality screen*), all DD-inclusive combinations of up to four features are considered. With 8 remaining candidates, this gives $\binom{8}{0} + \binom{8}{1} + \binom{8}{2} + \binom{8}{3} = 1 + 8 + 28 + 56 = 93$ combinations (DD alone, DD paired with one other, DD with two others, DD with three others). Each is estimated with a single Gibbs sampler seed, and combinations are eliminated if the HMM fails basic sanity checks: the two states must separate meaningfully, and the model must correctly classify known crises (the dot-com bust, the GFC, and COVID) as panic while keeping non-crisis periods calm.¹⁰

In the second pass (*portfolio ranking*), surviving combinations are evaluated through the full cross-sectional pipeline: HMM to XGBoost to long-short portfolio. This answers the question that matters most: does a better regime signal translate into better stock selection? Combinations are ranked by out-of-sample Sharpe ratio. The top 10 advance.

In the third pass (*stability validation*), the finalists are tested for robustness across random seed partitions. A combination that achieves a high Sharpe on one lucky seed combination but collapses on another is unreliable. We run 20 independent experiments with different seed draws and classify combinations as stable ($\sigma < 0.10$), moderate ($0.10 \leq \sigma < 0.15$), or unstable ($\sigma \geq 0.15$).

Selected features The four-feature combination DD+DISP+REL_N+CS ranks first on both ensemble Sharpe and stability mean (Table 2). The Sharpe ratios in the table are based on the selection pipeline’s seed budget (10 HMM seeds, 20 XGBoost seeds); the production configuration (200 HMM seeds, 50 XGBoost seeds) raises the Sharpe to 1.11 ($\sigma = 0.08$; Section 5) because additional seed averaging reduces MCMC and model noise.

¹⁰Specific thresholds: ESS > 50, regime separation significant for at least $\min(2, D)$ features, average $\pi_t^{\text{filter}} > 0.5$ during GFC and COVID, < 0.4 during 2003–2006 and 2013–2019, > 0.3 during the dot-com crash.

D	Combination	Sharpe	Mean	σ
4	DD+DISP+REL_N+CS	0.91	0.84	0.11
4	DD+DISP+CS+TERM	0.82	0.55	0.15
4	DD+DISP+REL_N+SKEW	0.80	0.69	0.08
4	DD+VOL+DISP+REL_N	0.80	0.77	0.09
3	DD+DISP+REL_N	0.77	0.77	0.10
4	DD+VOL+DISP+TERM	0.71	0.67	0.16
3	DD+VOL+CS	0.69	0.43	0.13
1	DD	0.68	0.67	0.14
4	DD+DISP+REL_N+TERM	0.62	0.45	0.18
3	DD+ADR+TERM	0.57	0.69	0.07

Table 2: Top 10 feature combinations from the three-pass selection pipeline (long-short, 10 bps costs). D = number of HMM features. Sharpe = ensemble Sharpe ratio (10 HMM seeds, 20 XGBoost seeds). Mean and σ are the mean and standard deviation of the Sharpe ratio across 5 independent experiments with different seed partitions, measuring robustness to random initialisation. DD+DISP+REL_N+CS ranks first on both ensemble Sharpe and stability mean.

Table 3 confirms the individual contribution of each feature. DD alone achieves the highest standalone Sharpe, but removing DD or REL_N from the full model causes the largest drops (-0.46 and -0.54). DISP and CS contribute marginally in isolation but reduce the Sharpe ratio’s standard deviation across seed partitions from 0.12 (DD alone) to 0.08 (full model), providing independent information from cross-sectional disagreement and corporate credit conditions that stabilises the regime classification.

HMM Input Features	XGB Sharpe	Δ vs Full
Full 4F baseline (DD+DISP+REL_N+CS)	0.97	—
<i>Include-one (single-feature HMM)</i>		
DD only (market drawdown)	0.66	-0.31
<i>Leave-one-out (three-feature HMM)</i>		
Drop DD	0.51	-0.46
Drop DISP	0.96	-0.01
Drop REL_N	0.43	-0.54
Drop CS	0.96	-0.01

Table 3: HMM input feature ablation (long-short, feature-selection seed counts). The absolute Sharpe levels differ from the production model (Table 5) due to fewer seeds; the relative comparisons (Δ) are the relevant quantities. DD alone nearly matches the full model. Removing DD or REL_N causes the largest drops (-0.46 and -0.54), while DISP and CS contribute marginally but improve seed stability (Section 4.2).

A note on feature selection and out-of-sample discipline The second pass of the feature selection procedure evaluates candidate feature sets using out-of-sample portfolio

performance, which means the selected combination (DD+DISP+REL_N+CS) partly reflects its test-period Sharpe ratio. Although no model parameters are re-estimated on the test set, the model *specification* (which features to include) is informed by test-period outcomes. This introduces a mild form of specification search. Three factors mitigate this concern. First, the quality screen in Pass 1 eliminates most combinations on purely in-sample criteria (regime separation, crisis classification) before any test-period evaluation. Second, the final selection in Pass 3 is based on cross-seed stability, not peak Sharpe: DD+DISP+REL_N+CS is chosen because it has the lowest standard deviation across seed partitions, not because it has the highest mean Sharpe. Third, and most importantly, the result is not unique to the selected combination. Table 2 shows that all feature counts from $D = 1$ to $D = 4$ produce positive long-short Sharpe ratios (0.57 to 0.91), and Table 3 confirms that even single-feature specifications (DD alone, Sharpe 0.68) and leave-one-out variants produce qualitatively similar results. The alternative train/test splits in Appendix G.1 further demonstrate that the regime-momentum interaction holds across different sample periods. Nevertheless, the reported Sharpe ratio of 1.11 should be interpreted with the caveat that the HMM feature set was selected with reference to test-period performance.

4.3 Regime Estimation Results

The convergence diagnostics described in Section 3.1.5 confirm that the Gibbs sampler with the selected features has converged. All posterior parameters exceed an ESS of 100, meaning the post-burn-in draws contain at least 100 effectively independent samples per parameter. The Gelman–Rubin $\hat{R}$ statistic is below 1.01 for every parameter across five independent chains (Appendix I.1), indicating that chains started from different initializations have converged to the same posterior distribution.

Table 4 and Figure 4 summarize the HMM output. All four features exhibit statistically significant separation between calm and panic regimes (95% credible intervals excluding zero), and the resulting signal correctly identifies all major stress episodes while remaining near zero during calm expansions.

	$\hat{\mu}_{\text{calm}}$	$\hat{\mu}_{\text{panic}}$	Δ	95% CI
DD_z	0.531	-0.976	-1.507	$[-1.794, -1.215]$
DISP_z	-0.423	+0.689	+1.113	$[+0.869, +1.353]$
REL_N_z	0.525	-0.895	-1.419	$[-1.591, -1.231]$
CS_z	-0.697	+0.203	+0.900	$[+0.626, +1.159]$

Table 4: HMM posterior regime-mean separation.

Notes: This table reports posterior mean estimates of regime-conditional feature means from the Bayesian two-state HMM, estimated via Gibbs sampling on the 1990–2010 training period.

All features are standardised to zero mean and unit variance. $\Delta = \hat{\mu}_{\text{panic}} - \hat{\mu}_{\text{calm}}$ is the separation between regimes. The 95% CI is the Bayesian posterior credible interval for Δ ; zero lies outside all intervals, confirming significant regime separation. DD = market drawdown; DISP = log cross-sectional return dispersion; REL_N = relative market participation; CS = credit spread (BAA–AAA).




Figure 4: Filtered regime probability π_t^{filter} (1990–2025). The dashed vertical line marks the train/test split (January 2011). Shaded regions indicate months where $\pi_t^{\text{filter}} \geq 0.5$ (panic classification).

4.4 Cross-Sectional Features

The cross-sectional models rank stocks each month using momentum signals and the filtered panic probability π_t^{filter} .

Momentum lookbacks Twelve trailing momentum signals are computed for each stock, corresponding to lookback horizons of 1 to 12 months:

$$\text{mom}_{k,t}^s = \prod_{\ell=1}^k (1 + r_{t-\ell}^s) - 1, \quad k = 1, \dots, 12, \quad (28)$$

where $r_{t-\ell}^s$ is the adjusted return of stock s in month $t - \ell$. The current month’s return is excluded from all lookbacks to avoid mechanical correlation with the forward return being predicted. Including all twelve horizons rather than a single lookback allows the models to learn how the full momentum term structure interacts with the regime signal.

Regime signal The filtered panic probability π_t^{filter} from the HMM enters as a single feature. Because it varies only at the monthly market level (every stock in the same month receives the same value), it functions as a conditioning variable rather than a stock-level predictor.

With the data and features defined, the next section presents the out-of-sample results of the three portfolio construction methods.

5 Results

All strategies are evaluated out-of-sample on 2011–2025 (167 months), using the long-short construction described in Section 3.2. The central question is whether momentum’s cross-sectional structure is regime-dependent, and whether exploiting this dependence requires nonlinear modelling. Section 5.1 compares the four model variants (DET, LR, Ridge, XGB) against unconditional momentum benchmarks; XGB emerges as the method that captures the regime-momentum signal best, and the rest of the chapter takes XGB as the chosen model and analyses how it does so.

5.1 Portfolio Performance

5.1.1 Results Overview

Table 5 reports the out-of-sample performance of all strategies.¹¹

	Ann. Ret	Ann. Vol	Sharpe	Max DD	β	NW t	Final \$
Market	11.9%	14.6%	0.84	-24.7%	1.00	–	4.8
Fixed 12-mo mom	-1.9%	26.1%	0.06	-72.2%	-0.66	0.24	0.8
Fixed 1-mo mom	3.1%	17.6%	0.26	-30.2%	-0.34	1.08	1.5
DET: Formula	-0.2%	23.6%	0.11	-63.8%	-0.59	0.43	1.0
LR	-1.3%	15.2%	-0.01	-42.4%	0.05	-0.04	0.8
XGB (mom+ π)	21.7%	19.5%	1.11	-22.8%	0.45	4.37***	15.3

Table 5: Out-of-sample portfolio performance.

Notes: Long-short (top/bottom decile, NYSE breakpoints), value-weighted, net of 10 bps one-way costs. β is the market beta. NW t uses Newey–West standard errors (6 lags). Test period: January 2011 to November 2025 (167 months). * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$.

Momentum’s cross-sectional predictability is regime-dependent, and exploiting this dependence requires both a probabilistic regime signal and a nonlinear model. XGB, an XGBoost regressor trained on the 12 momentum horizons together with the HMM filtered probability π_t^{filter} , achieves an annualised Sharpe of 1.11 with a six-factor alpha of 24.1% ($t = 4.81$; Appendix G.3). Block bootstrap paired tests reject equality of Sharpe ratios between XGB and each benchmark in Table 5 at $p < 0.025$; XGB’s own 95% Sharpe confidence interval is $[0.67, 1.53]$, comfortably above zero, while every benchmark’s interval straddles zero (Appendix G.4). An expanding-window backtest extending the OOS sample to 30 years (1995–2024) reproduces this Sharpe within sampling noise on the 167-month overlap with the main test period (0.88; see Section 5.3.3, which also covers the dot-com bust and GFC out-of-sample). Every other specification fails to produce economically meaningful long-short performance, and the pattern of failures is informative.

¹¹Sharpe convention details (raw vs. excess returns; arithmetic-vs-geometric mean) are in Appendix B.3.

Unconditional 12-month momentum achieves a Sharpe of 0.06 with a maximum draw-down of -72.2% , exposing the crash vulnerability documented by [Daniel and Moskowitz \(2016\)](#) and [Barroso and Santa-Clara \(2015\)](#). A deterministic regime switch (DET) improves this marginally to 0.11. The logistic regression (LR), given identical features to XGB, collapses to -0.01 ; a Ridge regression with the same 13 features but a continuous-return target (matching XGBoost’s) fails more decisively, with Sharpe -0.57 and a coefficient on π_t^{filter} of $+0.003$ (Appendix G.8). The failure is therefore linearity itself, not the classification target. Hand-engineering the missing nonlinearity into LR through π^2 and $\pi \times \text{mom}_k$ terms lifts it to 0.32, establishing that the regime-momentum interaction is real while also showing that pre-specified interactions recover less than a third of what XGBoost finds automatically. Adding firm-level fundamentals to XGB’s feature set degrades performance rather than improving it (Section 5.1.3), evidence that π_t^{filter} ’s role is specifically to condition the momentum term structure.




Figure 5: Cumulative wealth paths (log scale) starting from \$1 in January 2011. Shaded bands indicate panic months ($\pi_t^{\text{filter}} \geq 0.5$).

5.1.2 Regime Signal Ablation

The regime signal is essential. Removing π_t^{filter} and retraining on the 12 momentum features alone drops the Sharpe to 0.43 (Table 6). Replacing π_t^{filter} with the four raw stress indicators (DD, DISP, REL_N, CS) performs even worse (Sharpe 0.26). With four additional dimensions the model must reconstruct the regime structure at each split, while the HMM distils those indicators into a single dimension that one tree split can condition on. The HMM’s role is not to provide a convenient feature but to solve regime

identification *before* stock selection, freeing the cross-sectional model to use the regime signal rather than construct it. If the model cannot extract a meaningful regime signal even when given the four raw indicators directly, it certainly cannot do so from momentum features alone, which are one further step removed from the underlying regime dynamics.

Regime Signal Variant	Ann. Ret	Ann. Vol	Sharpe	Max DD
XGB + π_t^{filter} (HMM)	21.7%	19.5%	1.107	−22.8%
XGB + raw indicators (DD, DISP, REL_N, CS)	3.2%	19.1%	0.257	−45.4%
XGB (no regime signal)	7.0%	20.4%	0.429	−41.0%

Table 6: Regime signal ablation (50-seed ensemble). The HMM signal outperforms both raw stress indicators and no regime signal.

5.1.3 Feature-Set Ablation

The regime signal ablation removes π_t^{filter} from the production model. The complementary question is whether *adding* widely used firm-level fundamentals strengthens the strategy. Ten characteristics are tested (seven value-quality-size measures and three cash-flow-based earnings-quality measures, real-time and four-month-lagged from Compustat; definitions in Appendix K); each variant is trained at production settings.

	Feat.	Sharpe	Calm	Panic	β	NW t (exc.)	Fund SHAP
Baseline (mom + π)	13	1.11	0.84	1.53	0.45	1.83*	0%
+ 10 fundamentals	23	0.89	0.69	1.18	0.07	0.61	44%
Fundamentals only	10	0.59	0.86	0.38	0.19	−0.49	100%
Fundamentals + π	11	0.64	0.59	0.72	0.03	−0.52	60%
Momentum + fund. (no π)	22	0.63	0.68	0.57	0.16	−0.29	54%

Table 7: Feature-set ablation (50-seed XGBoost ensembles, production hyperparameters). Full Sharpe, calm-period Sharpe, and panic-period Sharpe are reported alongside market beta and the Newey-West t -statistic on excess returns over the market (6 lags), matching the convention of Table 5. “Fund SHAP” is the share of total SHAP attention absorbed by the ten fundamental features.

Adding the fundamentals to the baseline reduces Sharpe from 1.11 to 0.89 (Table 7). The new features absorb 44% of total SHAP, redirecting attention away from the momentum term structure (54% to 27%) and the regime signal (46% to 29%) without net improvement. Beta falls from 0.45 to 0.07 and maximum drawdown tightens from −22.8% to −18.3%, so the augmented variant behaves more like a low-beta market-neutral strategy than a high-alpha regime-conditioned one.

The fund-only variant identifies the source of the dilution. Its calm Sharpe of 0.86 essentially matches the baseline’s 0.84, but its panic Sharpe of 0.38 is a fraction of the baseline’s 1.53: stable firm attributes do not flip direction when the market enters panic.

Adding the regime signal to fundamentals ($\text{fund} + \pi$) lifts panic Sharpe to 0.72 but remains well short of the full model, because without the momentum term structure there is no directional relationship for the regime signal to reorganise. The regime-momentum interaction is the load-bearing mechanism; fundamentals cannot substitute for it. A regime-conditional feature-set design is discussed in Section 6.3.

5.2 How the Model Reorganises Momentum Across Regimes

Section 5.1 established that XGB is the only specification that produces meaningful out-of-sample long-short returns. This subsection asks what XGB learned. The answer is not that the model simply puts more weight on short-term momentum during stress or more weight on long-term momentum during calm periods. Instead, the model learns different cross-sectional momentum shapes. In calm markets, the long leg resembles a conventional momentum portfolio, selecting stocks with above-average intermediate-to-long horizon returns. In acute panic, the long leg reverses, selecting stocks whose past returns are below the cross-sectional average across the entire 1–12 month term structure. The remainder of this subsection documents that reorganisation in three steps. SHAP establishes that both the regime signal and the momentum horizons are decision-relevant (Section 5.2.1). Cross-sectional z-scores then show *what shape* the model actually buys (Section 5.2.2). Clustering reduces the noisy monthly heatmap into four recurring selection modes, two continuation and two reversal.

5.2.1 Per-Horizon Momentum SHAP

XGB’s predictions are decomposed using SHAP values (Lundberg and Lee, 2017), which measure how much each input contributes to the model’s prediction (definitions in Appendix C.1). Aggregate feature importance is split almost evenly. π_t^{filter} accounts for 46% of total SHAP, the 12 momentum horizons together for 54% (Table 8). The near-even split supports a context-variable role for π_t^{filter} . It is neither a minor control nor a substitute for momentum, but a routing variable that determines how the model uses the momentum term structure. Between regimes, momentum’s share rises from 50% in calm to 58% in panic while π_t^{filter} ’s share falls from 50% to 42%, with the model leaning on momentum more heavily during stress.

Feature	Overall	Calm	Panic
Momentum (12 horizons)	54%	51%	59%
π_t^{filter}	46%	49%	41%

Table 8: SHAP feature importance shares for XGB. Each column sums to 100%.

Beyond the aggregate split, the model’s reliance on each horizon shifts between regimes. In the calm long leg, months 8–12 together account for 66.0% of long-leg momentum

SHAP, with months 11 and 12 the most influential individual horizons (18.6% and 16.8%). In the panic long leg, months 7–12 account for 67.2%; the same intermediate-to-long band carries roughly two-thirds of the prediction signal in both regimes, mildly shifted toward the shorter end during panic. The sharpest within-regime move is on the panic short side, where month-1 SHAP rises to 10.6% versus 5.3% in calm, so the model becomes more sensitive to recent returns when identifying stocks to short during stress. SHAP tells us *which* horizons are decision-relevant; it does not tell us whether the model rewards or penalises high values at those horizons. A horizon can matter because the model buys winners at it, buys losers, or uses it only conditional on other horizons. To identify the economic content of the predictions, the next subsection turns from feature attribution to the cross-sectional characteristics of the stocks actually selected into the long leg.

5.2.2 Clustering Selection Modes Across Market States

The regime-conditional Sharpe is 1.53 in panic against 0.84 in calm (Table 9). Where exposure-scaling methods reduce positions during stress to limit losses, XGB instead reconstructs the portfolio (Section 5.3.2); the rest of this subsection shows that the regime signal flips the way the model uses the momentum term structure, from selecting winners in calm to selecting laggards in panic.

	Full	Calm	Panic
Market	0.84	0.64	1.13
Fixed 12-mo mom	0.06	0.08	0.05
Fixed 1-mo mom	0.26	0.52	−0.01
DET: Formula	0.11	0.18	−0.01
LR	−0.01	−0.29	0.39
XGB (mom+ π)	1.11	0.84	1.53

Table 9: Regime-conditional Sharpe ratios (108 calm months, 59 panic months).

To trace where in the cross-section the model picks, we use cross-sectional z-scores: for each stock in the long or short leg, the z-score at horizon k is $(mom_k^s - m\bar{o}m_k)/\sigma_k$, computed from the cross-sectional mean and standard deviation of all stocks that month. Figure 6 plots the long-leg z-curve month by month.




Figure 6: Long-leg cross-sectional z-scores by month and momentum horizon, 2011–2025. Each row is one of 167 test months sorted chronologically; each column is one of the 12 momentum horizons. Color shows the deviation of the long leg’s mean momentum at horizon h from the cross-section average that month.

Selection is not uniform across months, but the heatmap shows the model does not switch randomly either: it repeatedly returns to a small number of recognisable selection modes. To summarise these modes, we cluster the 167 monthly z-score profiles with K -means ($K = 4$, treating each month as a 12-dimensional point). The partition is highly stable across random initialisations.¹² The four groups nest inside a coarser two-fold partition: clusters 1 and 2 are the 115 months in which the long leg sits above the cross-section average (“above-average baskets”, $\bar{\pi} = 0.28$), and clusters 3 and 4 are the 52 months below (“below-average baskets”, $\bar{\pi} = 0.52$).

Per-cluster descriptive statistics are in Table 10; the full feature signature (cross-section context, raw long-leg pick momentum at each horizon, long-leg z-curve at each

¹²The same labels emerge across six independent random seeds for 163 of 167 months (97.6%), with mean pairwise Adjusted Rand Index 0.96. Pairwise centroid distances in 12-dimensional z-space range from 1.37 (clusters 1 and 2, both above-average) to 4.90 (clusters 1 and 4, the extremes), against a typical within-cluster scatter of 0.6 to 0.9.

horizon) is in Appendix Table 27. Each cluster’s z-curve panel appears immediately above its description (Figures 7–10); thin lines are member-month z-curves and the bold line is the cluster centroid.

Feature	1 (calm)	2 (mild)	3 (post-panic)	4 (deep crisis)
Months (n)	53	62	31	21
$\bar{\pi}_t^{\text{filter}}$ (current panic prob.)	0.21	0.34	0.32	0.81
π 6-mo prior frequency	0.27	0.37	0.30	0.51
π 12-mo prior frequency	0.31	0.32	0.41	0.30
Cross-section avg stock return	+14%	+7%	+13%	−9%
Cross-section std of 9–12 mo momentum	0.60	0.59	0.85	0.45
Cross-section skewness of 1–4 mo momentum	4.6	7.2	6.4	3.7
Strategy trailing 12-mo Sharpe	1.22	1.04	1.15	1.08
Long-leg avg z-score $\bar{z}$	+0.57	+0.19	−0.29	−0.82
Cluster Sharpe	0.92	0.93	1.21	1.82

Table 10: $K = 4$ cluster descriptors, 2011–2025 test period, $n_{\text{total}} = 167$ months. Each column is one cluster; rows are grouped (top to bottom): cluster size; regime context (current π_t^{filter} and prior 6-/12-month panic frequencies); cross-section state (average-stock trailing return; cross-stock standard deviation of 9–12 month momentum; cross-stock skewness of 1–4 month momentum); the strategy’s trailing 12-month Sharpe entering the cluster’s months; long-leg average z-score $\bar{z}$ (the model’s picks, averaged across the 12 momentum horizons); and the cluster Sharpe (annualised, block bootstrap with block size 6, 5,000 reps). The 9–12 month and 1–4 month windows are chosen as the ones with the largest cross-cluster spread for each measure.




Figure 7: Cluster 1 (calm continuation): z-curves of the 53 member months and centroid. The centroid is humped and uniformly positive, peaking at +0.80 at month 8.

Cluster 1 (calm continuation): a quiet, broadly trending market. With current π_t^{filter} at 0.21 (the lowest of any cluster), the model has high confidence the regime is calm and commits to its standard winner-selection mode. The long-leg z-curve is humped, with picks only mildly above average at short horizons ($z \approx +0.26$ at month 1) but clearly differentiated at intermediate horizons (peak $+0.80$ at month 8). In absolute terms, raw pick momentum reaches $+60\%$ at month 9, reproducing the buy-side of [Jegadeesh and Titman \(1993\)](#)’s twelve-minus-one strategy. The intermediate-horizon peak matches [Novy-Marx \(2012\)](#)’s documented dominance, and the underlying mechanism is [Hong and Stein \(1999\)](#)’s gradual information diffusion, in which news enters prices slowly so winners keep winning until the information is fully absorbed.

These 53 months are a stable, mature bull market, the conditions under which classical momentum works best. The cross-section is positive at every horizon (average stock up $3\%/13\%/23\%$ over 1/6/12 months), so there is a coherent up-trend for the model to ride. The rise is broad rather than driven by a few outliers (cross-section skewness $4.6/6.2/6.7$, below the test-period average of $5.8/6.6/7.1$), which means winner-selection picks up genuine widespread participation rather than chasing a handful of unrepresentative names. Recent panic activity is at its lowest (27% of the prior 6 months in panic), so the cluster sits far from the post-bear-market rebounds where unconditional momentum is known to crash, and continuation can pay safely. Consistent with this favourable setup, the strategy enters these months with the highest trailing 12-month Sharpe of any cluster (1.22). The cluster is concentrated in 2013–2014 and runs in stretches of up to 21 consecutive months, the kind of extended UP-market state in which [Cooper et al. \(2004\)](#) document that momentum profits are most concentrated.

Cluster 1 earns a Sharpe of 0.92, the lowest of the four. The Sharpe is moderate because the trend is mature; continuation still pays, but the strategy is riding an established rally rather than catching a turning point. Cluster 1 is a sanity check on the architecture, since the model recovers classical momentum selection here, confirming the regime-conditioning architecture adds new behaviour during stress without overriding the unconditional pattern that already works in calm.




Figure 8: Cluster 2 (mild continuation): z-curves of the 62 member months and centroid. The centroid is mildly positive throughout, peaking at +0.32 at month 8.

Cluster 2 (mild continuation): a choppy market with no broad direction. With current π_t^{filter} in an ambiguous middle band (0.34, neither clearly calm nor clearly panic), the model cannot commit to either winner-selection or laggard-selection. The long-leg z-curve has the same humped shape as cluster 1 but at a much lower amplitude (z-mean +0.19 against the test-period long-leg average of +0.09, peak +0.32 at month 8); the model picks up the same intermediate-horizon signal, but it is weaker here. Performance comes from occasional outlier capture rather than systematic selection.

These 62 months are an ambiguous transition state, with recent stress still lingering but no clear up-trend in place. The 6-month panic frequency is 0.37, above the test-period average of 0.34, while the overall cross-section trend is mild ($\text{mom}_{\text{overall}} = 0.07$, below the test-period average of 0.08). The cross-section is only mildly positive on average (+2%/+8%/+11% over 1/6/12 months), and cross-section skewness is the highest of any cluster (7.2/7.6/8.1), with a small number of extreme winners pulling the mean above zero while the median stock's return stays near flat, so even the modest positive level is not broad-based. The strategy reflects the same ambiguity, entering with the lowest trailing 12-month Sharpe of any cluster (1.04). Cluster 2 is the largest of the four ($n = 62$) but appears only in short bursts (max 4 consecutive months), consistent with its role as a brief transition between regimes rather than a stable state. The muted shape is consistent with [Cooper et al. \(2004\)](#)'s finding that momentum profits depend on the prior market state.

Cluster 2 earns a Sharpe of 0.93, statistically indistinguishable from cluster 1, and exposes the limit of the regime-momentum mechanism. When the cross-section carries no coherent direction, the model cannot build a high-conviction long leg even with the

regime signal available. The strategy earns roughly flat returns in these months and gets its small edge through occasional outliers rather than systematic selection.




Figure 9: Cluster 3 (post-panic recovery): z-curves of the 31 member months and centroid. The centroid sits below the cross-section average at every horizon, between -0.22 and -0.36 .

Cluster 3 (post-panic recovery): a recovering market with high long-horizon dispersion. Moderate current π_t^{filter} (0.32) combined with the highest 12-month prior frequency (0.41) is the feature combination the model has learned to read as a recently-resolved panic. The trees route into a post-bear-rebound branch, building the long leg from cross-sectional laggards rather than winners. The long-leg z-curve sits below the cross-section average at every horizon (z-mean -0.29), with the deepest below-average at month 4 (-0.36) and month 12 (-0.35). The picks’ absolute trailing returns range from -6% to $+1\%$ while the broader cross-section is up double digits. The picks rotate more in this cluster than in any other: only about 68% of the long-leg names carry over from one cluster-3 month to the next consecutive cluster-3 month (the lowest carry-over fraction of any cluster), as some recovery candidates catch up and others remain behind.

These 31 months follow a panic-heavy year (41% of the prior 12 months were panic, the highest of any cluster) while current panic probability is only moderate ($\pi_t^{\text{filter}} = 0.32$), so the cross-section is recovering. The average stock is up $+1\%/+11\%/+27\%$ over 1/6/12 months, with cross-section dispersion the highest of any cluster at both the 5–8 and 9–12 month horizons (0.56 and 0.85), as some stocks rebound and others lag. Cluster 3 is heavy in 2021 (9 of 31 months, the post-COVID rotation) and 2023, both of which share the dynamic [Daniel and Moskowitz \(2016\)](#) identify as crash-prone for unconditional momentum, where past losers (high- β stocks) rally disproportionately and past winners (low- β defensives) lag, so the cross-sectional ranking inverts. Cluster-level

leg-betas (Appendix G.7, Table 28) confirm the inversion empirically: fixed 12-month momentum has the most negative long-minus-short β spread in cluster 3 (-0.75 , the deepest inversion of any cluster), while XGB has flipped the selection so its long-leg β exceeds its short-leg β ($+0.49$). The increasing dispersion across horizons also matches Lee and Swaminathan (2000)’s lifecycle interpretation, in which momentum signals evolve from continuation to reversal as information ages.

Cluster 3 earns a Sharpe of 1.21 on these laggard picks. The novelty is the cluster’s identification of a dispersed-recovery selection mode: while the cross-section is up double-digits and dispersion is at its highest, the long leg is built from stocks still below average at every horizon. This is qualitatively distinct from the short-horizon reversal of Jegadeesh (1990) (one-month) and Lehmann (1990) (weekly), which fire on a signal at a single short horizon, and from De Bondt and Thaler (1985)’s long-horizon contrarian reversal at 3–5 year windows; cluster 3 selects names that sit below the cross-section average *uniformly across the 1–12 month range*, a region of horizon space targeted by neither classical reversal.




Figure 10: Cluster 4 (deep crisis): z-curves of the 21 member months and centroid. The centroid sits uniformly between -0.73 and -0.88 across all 12 horizons.

Cluster 4 (deep crisis): an active panic with the whole market falling. The market has been falling for a sustained period. The probability of panic is at its highest ($\pi_t^{\text{filter}} = 0.81$), 51% of the prior 6 months were also panic (the highest of any cluster), and the cross-section is uniformly negative, with the average stock down 5%/11%/7% over 1/6/12 months. With π at this extreme, the model has unambiguous panic detection and commits fully to deep-crisis laggard-selection. In this falling market the model picks the most beaten-down names. The long-leg z-curve sits uniformly between -0.73 and -0.88 across all 12 horizons (deepest around month 9), meaning the picks are roughly

0.7 to 0.9 cross-section standard deviations below the average stock at every horizon. Their absolute returns deepen from -15% over the prior month to -48% over the prior 12 months. The strategy bets on a rebound, since stocks that have fallen the furthest in the panic are typically those with the highest market beta, and these high- β losers rally disproportionately once the stress resolves. This is the same high- β -rebound mechanism as cluster 3, but at the panic peak rather than during recovery.

Critically, only 30% of the prior 12 months were panic, in line with the test-period average; these are sharp, short panics rather than prolonged year-long deteriorations. The reversal bet rests on this premise; Section 5.3.2 examines what happens to the strategy when the panic is instead sustained. Cross-section dispersion is the lowest of any cluster at both the 5–8 and 9–12 month horizons (0.35 and 0.45), and 1–4 month cross-section skewness is also the lowest (3.67), so stocks are moving down together rather than fanning out. Eleven of 21 months are in 2022 (Fed-hiking cycle), with corroborating instances in April 2020 (COVID trough), October 2015 and January 2016 (China devaluation, oil collapse), and January 2023.

Cluster 4 earns the highest Sharpe of any cluster (1.82, 95% block-bootstrap CI $[0.98, 2.68]$), with a mean monthly return of $+3.13\%$. The novelty is twofold. First, this is a term-structure-wide reversal portfolio with no direct analogue in the reversal literature. Jegadeesh (1990) (one-month) and Lehmann (1990) (weekly) document short-term reversal; Nagel (2012) shows it spikes during crises as compensation for liquidity provision; De Bondt and Thaler (1985) document long-term contrarian reversal at 3–5 year horizons. Each targets a single horizon band, while cluster 4’s uniform loading spans all 12 horizons together. Second, where Daniel and Moskowitz (2016) document the beta-reversion mechanism as a *risk* to be scaled away from, XGB exploits it as a strategy. The regime signal makes the contrarian bet loadable, and it pays off precisely because these panics are sharp and short rather than prolonged. That π_t^{filter} ’s combined long-and-short SHAP share is *lowest* here (0.24) reflects the routing structure, since once in the deep-crisis branch, momentum splits alone produce the loser portfolio. Cluster-level leg-beta and SHAP-share diagnostics for the full $K = 4$ partition are in Appendix G.7.

5.2.3 Discussion

The $K = 4$ partition isolates the central economic finding of this thesis. XGB does not simply shorten the momentum lookback during stress, nor does it use π_t^{filter} as a constant return predictor. Instead, the model learns four recurring selection modes and lets the regime signal route between them. Clusters 1 and 2 are above-average baskets (continuation modes); clusters 3 and 4 are below-average baskets (reversal modes). The transition from continuation to reversal is not captured by any single fixed lookback or by any linear coefficient on a regime variable. It requires reading the same momentum horizons differently inside each branch of the regime-conditioned tree.

This is why the linear baselines on identical features collapse (Section 5.1). A linear model assigns one slope per momentum horizon, so it can answer “does high momentum predict high return?” but not “does high momentum predict high return in calm and low return in panic?”. XGB’s tree splits encode that state-dependent reversal directly. Once a tree has branched on π_t^{filter} , the same momentum input can be rewarded in one branch and penalised in the other. The regime signal is therefore not a return predictor in its own right but a routing variable that determines which momentum shape the model treats as attractive.

Economically, this reframes momentum’s crash-management problem as a stock-selection decision rather than an exposure-scaling one. Daniel and Moskowitz (2016) and Barroso and Santa-Clara (2015) reduce gross exposure when crash conditions arise; XGB keeps full exposure but reorganises which stocks fill each leg. The cluster-level CAPM betas confirm this empirically (Appendix G.7, Table 28). XGB maintains long-leg β above short-leg β in every cluster (spreads +0.68, +0.40, +0.49, +0.43 for clusters 1–4), while fixed 12-month momentum exhibits the D&M leg-beta inversion in every cluster (−0.24, −0.87, −0.75, −0.42). Where exposure-scaling avoids the crash mechanism by shrinking, XGB avoids it by re-ranking.

The regime-aware adaptation also concentrates on the reversal side. Clusters 3 and 4 differ by 0.61 in Sharpe while clusters 1 and 2 are indistinguishable (0.01); the model distinguishes reversal environments more sharply than trend-following ones. The cluster 4 reversal mechanism, which drives the largest single-cluster Sharpe (1.82), rests on a specific premise about the nature of panic-period declines, which Section 5.3.2 stress-tests.

Three contributions to the existing momentum literature (Cooper et al., 2004; Daniel and Moskowitz, 2016; Goulding et al., 2023) follow: (i) a term-structure-wide reversal portfolio in deep crisis with no direct analogue in the reversal literature; (ii) a regime signal whose marginal value scales with cross-sectional ambiguity rather than acting as a constant feature; (iii) selection by momentum *level* rather than momentum *rank*.

5.3 Strategy Pitfalls

The strategy’s headline performance comes with three pitfalls. First, the baseline volatility is high (around 19.5%) and some investors will need to dial it down; Section 5.3.1 traces the risk-return trade-off offered by a CRRA risk-aversion control. Second, the panic-regime reversal rests on the premise that price declines during stress are temporary; a sustained economic deterioration would expose the strategy to reversal failure, which Section 5.3.2 quantifies with a synthetic-recession stress test. Third, the 2011–2025 test sample does not include the worst historical stresses; Section 5.3.3 extends the evaluation to 30 years (1995–2024) via expanding-window backtests covering the dot-com bust, the GFC, and

the 2009 momentum crash.

5.3.1 Risk-Aversion Control

The first pitfall is volatility itself: an investor may prefer lower portfolio risk at the cost of some expected return. We incorporate a risk-aversion parameter γ into the strategy and trace the out-of-sample trade-off across $\gamma \in [0, 1]$.

The natural starting point is the mean-variance utility $r - \frac{\gamma}{2}\sigma^2$ (Markowitz, 1952), but it is the wrong lens here. Mechanically, applied as a scoring function the penalty is subtracted symmetrically from every stock’s score, pushing the long book toward low-volatility names and the short book toward high-volatility names: a betting-against-beta portfolio whose realised volatility rises rather than falls as γ grows (Appendix D confirms this across several variants). Economically, mean-variance penalises variance in *absolute* terms, but real investors behave as if risk aversion is *relative*, with a stable tolerance for proportional risk that does not change with wealth. Constant relative risk aversion captures this; mean-variance does not.

We therefore adopt a CRRA utility that penalises $\log \sigma$ instead of σ^2 , and evaluate each stock’s utility separately under a long direction and a short direction. The volatility penalty enters identically in both: a volatile stock is equally unattractive as a holding or as a short position, so the investor’s aversion to risk does not flip sign with trading direction. Longs are the top decile of long-utility across the full cross-section; shorts are the top decile of short-utility. The full specification is given in Appendix E.

Risk-Aversion Sweep: Sharpe, Return, Volatility, Drawdown vs. γ




Figure 11: Risk-aversion sweep. Sharpe ratio, annualised return, annualised volatility, and maximum drawdown as a function of γ over $[0, 1]$. Red markers highlight $\gamma \in \{0, 0.1, 0.5, 1.0\}$.

Figure 11 shows the resulting Pareto curve.¹³ At $\gamma = 0.1$, Sharpe stays at 1.01 while maximum drawdown tightens from -25% to -18% , a near-free reduction in tail risk. At $\gamma = 0.5$, annualised volatility halves from 19.7% to 11.2% , drawdown halves to -12% , and Sharpe settles at 0.78. At $\gamma = 1.0$, volatility falls to 5.2% (roughly a quarter of the market's) and Sharpe remains positive at 0.46. An investor can choose any point on this curve based on their volatility budget or drawdown tolerance. Beyond $\gamma \approx 1.5$, the risk penalty dominates the return signal and the long-short spread collapses into dollar-neutral noise. The framework does not blow up; it simply stops producing a directional bet.

5.3.2 Stress Test: Prolonged Bear Market

The reversal mechanism documented in Section 5.2.2 (cluster 4) rests on a specific premise: that price declines during panic are temporary (driven by liquidity pressure) rather than permanent (driven by economic deterioration). The historical stress periods in the test sample support this: XGB earned $+29.9\%$ during the COVID crash (February–April

¹³The $\gamma = 0$ baseline reflects the CRRA dual-utility implementation rather than the headline XGB scoring; small numerical differences from Section 5.1 (volatility 19.50% , maximum drawdown -22.8%) are within seed noise.

2020) and +21.5% during the 2022 bear market (January–October 2022). In both cases, the decline was temporary and reversed within months.

But the premise can fail. In a sustained economic deterioration (a depression or structural decline), beaten-down stocks fall because their businesses are failing, not because of temporary mispricing. We call this *reversal failure*: the regime classification is correct, but the momentum structure within that regime has changed. To quantify this vulnerability, we simulate recessions of increasing duration, replacing consecutive months with XGB’s average losing panic-month return (-3.44% per month) at every possible insertion point.

Recession Duration	Market Loss	XGB Loss	MDD (worst case)
Baseline (no recession)	—	—	-23%
3 months	-10%	-10%	-38%
6 months	-19%	-19%	-42%
12 months	-34%	-34%	-49%
18 months	-47%	-47%	-61%
24 months	-57%	-57%	-69%

Table 11: Prolonged bear market simulation. Each row injects a recession during which XGB loses 3.44% per month (its average losing panic-month return). MDD is the worst-case maximum drawdown across all possible insertion points.

The strategy can absorb substantial downturns. A 12-month recession produces a worst-case drawdown of -49% , after which calm-period continuation resumes. Only a depression-length downturn of 24 months produces a drawdown of -69% . This mirrors the vulnerability identified by [Daniel and Moskowitz \(2016\)](#) for traditional momentum, but on the opposite side: standard momentum crashes when winners stop winning; the regime-conditioned strategy would struggle when losers stop reverting.

[Daniel and Moskowitz \(2016\)](#)’s own response to momentum’s crash risk is exposure scaling. The two approaches use regime information differently: D&M scale *exposure* (how much to hold), XGB changes *which stocks* are held. On the same data and transaction costs, XGB (Sharpe 1.11, CAPM alpha 23.4%) outperforms D&M managed momentum (Sharpe 0.49, CAPM alpha 10.3%). A hybrid that adds D&M-style scaling on top of XGB as a circuit breaker (Section 6.3) does not improve performance in the current test period because no prolonged downturn occurred, but would help most in cluster 4, where the long leg sits term-structure-wide negative and the reversal bet is most exposed to failure. Section 5.3.3 provides direct historical evidence on how the strategy actually behaved during the dot-com bust and the GFC, complementing this synthetic stress test.

5.3.3 Historical Stress: Out-of-Sample Evidence

The 14-year test sample used so far (2011–2025) misses the worst momentum-relevant stress events: the dot-com bust, the global financial crisis, and the 2009 momentum crash documented by [Daniel and Moskowitz \(2016\)](#). To extend the evaluation, the full pipeline (HMM + XGBoost) is re-estimated on an annual expanding-window schedule: for each year T from 1995 onward, the model is trained only on data available through 31 December of year $T-1$ and then used to make predictions for year T . Stitching these year-by-year predictions together produces a 30-year out-of-sample monthly return series spanning 1995–2024. The series is genuinely out-of-sample at every point: no future information enters any year’s prediction. Methodology details are in Appendix [G.2](#).

Period	N months	Sharpe	Cumulative	Max DD
Full sample, 1995–2024	359	0.50	1,405.5%	-64.8%
1995–1999	60	0.50	50.8%	-25.5%
2000–2002 dot-com	36	-0.15	-35.9%	-64.8%
2003–2006 bull	48	0.70	37.5%	-12.7%
2007–2009 GFC	36	0.26	11.1%	-36.3%
2010–2019 calm	120	0.67	155.4%	-17.9%
2020–2024 COVID era	59	1.23	299.1%	-22.2%

Table 12: Sub-period decomposition of the 30-year expanding-window out-of-sample back-test. Sub-period rows aggregate by calendar year. The 2011–2024 sub-window of this test (167 months, the same window as the main test in Section [5.1](#)) reproduces the production Sharpe within sampling noise: 0.88 vs production 1.11.

Worst observed stress: the dot-com bust The 2000–2002 dot-com bust produced the strategy’s largest historical drawdown of -64.8% peak-to-trough. Restricted to the bear-market proper (March 2000–October 2002), the strategy lost -43.8% cumulative over 32 months (Sharpe -0.39). This is the historical realisation of the reversal-failure vulnerability described in Section [6.2](#): a prolonged sector-rotating bear in which beaten-down stocks did not promptly recover, so the panic-regime reversal selection bled. The strategy survived the episode and resumed positive returns thereafter (the 2003–2006 sub-sample produces Sharpe 0.70), but the realised drawdown is severe.

Stress at the targeted failure mode: the 2009 momentum crash The single most informative episode is March–December 2009, the period [Daniel and Moskowitz \(2016\)](#) document as a severe momentum crash. Under the same long-short construction used throughout this thesis (NYSE breakpoints, value-weighted, 10 bps cost), unconditional 12-month momentum lost -62.6% cumulative during those ten months. Under expanding-window out-of-sample, the regime-conditional strategy gained $+20.5\%$ cumulative (Sharpe 0.73) during the same window. Across the full GFC episode (October 2007–June 2009)

the strategy returned +22.8% cumulative (Sharpe 0.48). This is direct historical evidence that the regime mechanism survives the specific stress event that defeats unconditional momentum.

Synthetic and historical envelopes agree The synthetic stress test’s worst-case bound (−69% drawdown at a 24-month synthetic recession; Section 5.3.2) contains the realised dot-com drawdown of −64.8%, so the simulation does not understate the historical risk.

Long-horizon survival Aggregating across all 30 years and all stress events, the strategy produces +1,405% cumulative return with a Sharpe of 0.50. Each historical drawdown is followed by recovery; no stress episode in the sample is unrecoverable. Sub-period samples are small enough that point estimates carry wide confidence intervals, so the qualitative pattern across periods is more informative than the levels.

5.4 International Robustness

To test whether the cross-sectional regime-momentum mechanism extends beyond US data, we replicate the full pipeline on UK and Japanese equities (Compustat Global, 1990–2025).¹⁴

Table 13: International cross-sectional model performance, 2011–2025 out-of-sample test period. XGB uses the US specification (12 momentum horizons, 50-seed XGBoost ensemble, decile-spread long-short construction, value-weighted by market equity) with locally-fitted HMM (200 production seeds, DD+VOL+REL_N feature set selected per region). Long and short legs are formed from unconditional decile breakpoints across all listed stocks (no NYSE-equivalent tier exists for either market). Net of 10 bps one-way transaction cost.

Strategy	UK			JP		
	Sharpe	Ann. Ret	Max DD	Sharpe	Ann. Ret	Max DD
Local market	0.62	6.7%	−25.2%	0.86	12.3%	−21.5%
Fixed 12-mo mom	0.45	9.3%	−62.6%	0.07	−0.5%	−66.3%
XGB	0.84	16.0%	−21.4%	0.50	6.0%	−19.8%

Both regions support the mechanism. UK XGB attains a Sharpe of 0.84 over 2011–2025, well above unconditional 12-month momentum’s 0.45 and the local market’s 0.62. JP XGB attains 0.50, substantially above unconditional momentum’s 0.07 but below

¹⁴Each region’s HMM is fit on locally-selected stress features (DD+VOL+REL_N for both, with BANK_REL replacing the US-specific Moody’s BAA-AAA spread in the candidate pool); details in Appendix J. The cross-sectional model is unchanged from the US specification, except that long and short legs are formed from unconditional decile breakpoints across all listed stocks rather than NYSE-equivalent breakpoints (no analogous tier exists for either market).

the JP market's 0.86, consistent with the documented weakness of Japanese momentum (Asness et al., 2013). Maximum drawdown is reduced from -62.6% (fixed momentum) to -21.4% in UK, and from -66.3% to -19.8% in JP.

Taken together, the results show that momentum's cross-sectional predictability shifts between continuation and reversal across market regimes, and that exploiting this shift requires nonlinear conditioning on a regime signal. The signal's value lies not in predicting returns directly but in conditioning how the model jointly uses the full momentum term structure, a multi-dimensional interaction only tree-based models can learn.

6 Conclusion

The central finding of this thesis is that, in the sample studied, momentum’s cross-sectional predictability shifts with market conditions: which lookback horizons carry predictive power, and in which direction, depends on the prevailing market regime. The regime signal captures this shift not by predicting returns directly but by conditioning how the model jointly uses the full momentum term structure, a multi-dimensional interaction that only a nonlinear model can exploit. These findings complement the exposure-scaling approach of [Daniel and Moskowitz \(2016\)](#) and [Barroso and Santa-Clara \(2015\)](#) along two distinct channels. First, the HMM filtered probability provides a continuous, smoothed regime signal, a probabilistic alternative to the binary bear-market indicators used to gate exposure-scaling rules. Second, the cross-sectional model uses that signal to reorganise which stocks fill each leg rather than to scale gross exposure: the regime signal conditions the direction of momentum selection, and the leg-beta inversion that drives traditional momentum’s crashes is absent in the resulting portfolio ([Appendix G.5](#)).

6.1 Contributions

The three contributions outlined in [Section 1](#) are supported by the empirical evidence. The term-structure reorganisation is visible directly in z-score profiles and does not depend on the modelling framework. The context-variable mechanism is confirmed by SHAP analysis showing that π_t^{filter} conditions the joint use of momentum horizons rather than predicting returns directly. The nonlinearity requirement is established by the failure of both classification-based and regression-based linear models on identical features. Together, these findings suggest that the fragility of momentum documented in prior work may partly reflect unconditional models averaging over a regime-dependent structure, rather than an intrinsic deficiency of the anomaly itself.

6.2 Limitations

Economic vulnerability The strategy’s primary vulnerability is a prolonged economic deterioration where beaten-down stocks do not recover because their underlying businesses are failing rather than because of temporary mispricing. The stress test in [Section 5.3.2](#) quantifies this: a 12-month recession produces a worst-case drawdown of -49% , while a 24-month depression pushes the worst-case drawdown to -69% ([Table 11](#)). Calm-period continuation provides a floor, but a sustained downturn where reversal consistently fails would be painful.

Stress-tested vulnerabilities The strategy’s largest historical drawdown is the dot-com bust (-64.8% peak-to-trough; -43.8% cumulative over the 32-month bear-market

proper, March 2000–October 2002, under expanding-window OOS; see Section 5.3.3). This is the empirical realisation of the reversal-failure vulnerability described above and falls within the envelope estimated by the simulation in Section 5.3.2 (−69% worst-case drawdown for a 24-month synthetic recession). Worst-case scenarios beyond the historical envelope (a sustained multi-year decline without any reversal, e.g. a depression-style outcome) remain testable only via simulation.

Test period scope The test period (2011–2025, 167 months) does not contain a prolonged recession. The 30-year expanding-window backtest (Section 5.3.3) confirms that XGB would have survived the 2009 momentum crash, but the dot-com result there is the historical realisation of the reversal-failure vulnerability discussed above. Performance is strongest in 2021–2025 (Sharpe 1.70), a period with frequent regime transitions; in calm-dominated years the model earns modest returns from standard continuation. A permanently calm market would reduce the strategy to ordinary momentum.

Frozen HMM parameters and regime drift The HMM is estimated on 1990–2010 and held fixed during the test period. This assumes that the statistical properties of calm and panic regimes remain stationary. If the market undergoes structural change (a permanent increase in baseline volatility due to algorithmic trading or shifting monetary policy regimes, for example), the trained thresholds may become miscalibrated, with structurally calm periods classified as panic. The sustained elevation of π_t^{filter} in recent years may partly reflect such drift rather than genuine prolonged stress. The 30-year expanding-window backtest in Section 5.3.3 (in which the HMM is re-estimated annually using only data available at the time) reproduces the production result on the 2011–2024 overlap, indicating that the regime structure is reasonably stable; online re-estimation would still provide a stronger guarantee.

Memoryless regime conditioning XGBoost observes only the current month’s features and has no memory of the π_t^{filter} trajectory, so a fresh spike and a sustained elevation are treated identically despite their different implications for reversal. Section 6.3 sketches a sequential extension.

Implementation frictions The analysis assumes frictionless short selling and applies a uniform 10 basis point one-way transaction cost. Cost sensitivity (Appendix B.2) confirms XGB remains profitable at 50 bps, but the uniform cost understates the borrow fees and hard-to-borrow constraints that would apply most strongly to the short leg for small or less liquid stocks. Value-weighting and NYSE decile breakpoints tilt the portfolio toward large, liquid names, partially mitigating this concern.

Specification search The HMM feature selection procedure uses out-of-sample portfolio Sharpe in its second pass, meaning the selected combination (DD+DISP+REL_N+CS) is informed by test-period outcomes. Section 4.2 discusses the mitigating factors in full (the quality screen in Pass 1, the cross-seed stability criterion in Pass 3, and the qualitatively similar performance of leave-one-out specifications, Table 3); but the reported 1.11 Sharpe should be read with this caveat.

Sample scope and causal inference The sample is restricted to US equities. The interpretation that the regime signal conditions momentum’s term structure is supported by multiple lines of evidence (SHAP, tree paths, ablation, LR-XGB divergence) but is inferred from observational data and should be read as consistent evidence rather than definitive proof. The Newey–West t -statistic on XGB’s excess return over the market is 4.37^{***} , well above conventional thresholds, and the factor-model alphas in Appendix G.3 confirm this strength across all five specifications (CAPM, FF3, Carhart, FF5, FF6: every $t > 4$).

6.3 Future Work

Several directions could extend this framework.

Predictive regime signal The current regime signal π_t^{filter} estimates the probability that the market is in panic *now*, conditional on data through time t . Because the HMM defines transition probabilities between states, a natural extension is to predict the regime one step ahead: $P(s_{t+1} = 1 \mid \mathbf{z}_{1:t}) = \sum_k \pi_t^{\text{filter}}(k) \cdot P_{k,1}$. Feeding this predictive probability into the cross-sectional model would allow the portfolio to reposition before a regime shift rather than reacting to it. The expected gain is concentrated in transition months, where the current framework trades on a stale view of the regime in precisely the months when the portfolio is furthest from its regime-conditional optimum. A richer extension applies the same recursion to produce a k -step forecast, turning the regime signal into a short-horizon expected-state distribution that the cross-sectional model can condition on directly.

Alternative nonlinear models XGBoost is one of several nonlinear model classes, and the central mechanism (that π_t^{filter} gates the use of momentum horizons) should in principle be recoverable by any sufficiently expressive architecture. The Random Forest comparison in Appendix F.3 partially tests this: bagging still selects π_t^{filter} as a material feature, but the drop in Sharpe (1.08 to 0.73) and the fall in its SHAP share (45% to 18%) show that the boosted additive structure is what elevates the regime signal from a marginal predictor to a primary gating variable. The more informative test is whether the mechanism survives under genuinely different inductive biases. An attention-based model

would learn horizon-by-regime weights directly and produce an interpretable per-horizon attention profile that could be compared against the SHAP decomposition in Table 8; a well-regularised multilayer perceptron with explicit $\pi \times mom_k$ and $mom_k \times mom_j$ interaction inputs would isolate whether the dependence on boosting reflects sequential residual fitting or a shortage of data for richer architectures. Recovering the same mechanism across these alternatives would strengthen the claim that the regime-conditional shift is a property of the data rather than of XGBoost specifically; failure would force a narrower interpretation tied to tree-based models.

Sequential regime conditioning XGBoost observes only the current month’s features and treats all months with the same π_t^{filter} identically. Two months both at $\pi_t^{\text{filter}} = 0.75$ carry the same input regardless of whether the signal just spiked from 0.1 (a fresh crash, when mean reversion is most likely) or has been elevated for six consecutive months (prolonged stress, when the reversal mechanism may be breaking down). In practice, the spike is where the long-reversal premium is largest and the sustained elevation is where the strategy is most exposed to failed reversals. A sequential model could distinguish these cases: candidates include a recurrent network such as an LSTM, a transformer over a fixed window of regime features, or a simpler feature-engineering approach that adds lagged π_t^{filter} values, a short-window rate of change, and a duration-since-entry counter to the existing input set. The HMM’s forward-backward recursion already integrates the full history into π_t^{filter} , but reducing that history to a single scalar discards exactly the information needed to distinguish a fresh spike from a prolonged one. Exposing the trajectory to the cross-sectional model, rather than only its filtered expectation, should improve decisions in the months where the current specification is weakest.

Regime-conditional feature sets Adding ten firm-level fundamentals to XGB reduces Sharpe from 1.11 to 0.89 (Section 5.1.3) because the new features absorb 44% of SHAP importance and dilute the regime-momentum signal. The feature-set ablation localises the dilution: fund-only attains a calm Sharpe of 0.86 (matching the baseline’s 0.84) but a panic Sharpe of only 0.38, because firm-level characteristics describe stable attributes that do not flip direction when the regime shifts. The natural conclusion is not that fundamentals are uninformative but that the optimal feature set may itself be regime-dependent. During panic, when reversal dominates, momentum and the regime signal carry the mechanism, and adding quality or valuation characteristics competes for splits without contributing to the primary gating rule. During calm, when the continuation premium is weaker and more diffuse, fundamentals may provide a useful second dimension. Several designs could test this. The simplest is a layered construction in which XGB produces the initial long-short ranking and fundamentals enter as a post-ranking filter that excludes bottom-quintile-quality names from the long leg and high-leverage names from the short leg. A more

ambitious variant trains two separate cross-sectional models, selected by π_t^{filter} , each with its own feature set. Both designs require careful out-of-sample validation because the regime-conditional split reduces effective sample size and is vulnerable to overfitting, but both offer a principled way to recover the information that fundamentals plausibly carry without diluting the mechanism that drives current performance.

Cross-validated feature selection The HMM feature-selection procedure used here (Section 4.2) evaluates candidate feature sets in part by their out-of-sample portfolio Sharpe, a step that draws on test-period outcomes and is acknowledged in Section 6.2. A more disciplined alternative would replace the test-period evaluation with a cross-validation procedure run entirely within the training partition, for instance a five-fold expanding-window cross-validation on the 1990–2010 panel that never touches the 2011–2025 test data, with each fold scoring candidate feature sets on its held-out validation slice. Rolling-window robustness checks at multiple training-end dates would additionally test whether the chosen feature set is stable across regimes or sensitive to the specific window the procedure observes. Adopting cross-validation as the primary selection method, with the in-sample feature ablation reported in Section 4.2 as a complementary stability check, would remove the test-period dependence from the feature choice and produce a more rigorous justification for the four-feature input.

Systematic hyperparameter dissection The XGBoost specification used throughout (depth 4, learning rate 0.05, 500 trees) was chosen via a test-period sweep over a small grid (Appendix F.3) and is therefore subject to the same specification-search caveat as the HMM features. A natural extension is a fine-resolution sweep of the hyperparameter space (tree depth, learning rate, ensemble size, regularisation strength, and column-subsampling rate), with each configuration evaluated not only on out-of-sample Sharpe but also on the cross-sectional z-scores of long- and short-leg selections across regimes. Because the central mechanism is a regime-conditional shift in the momentum term structure rather than headline Sharpe alone, the diagnostic question is whether the cross-sectional z-score reorganisation documented in Section 5.2.2 survives changes in model capacity, or whether the deep-crisis reversal selection pattern is specific to the production specification. A rolling-window cross-validation across multiple training-end dates would further test stability of the optimum across market regimes. The expected outcome is a tight cluster of high-Sharpe configurations with similar z-score profiles, in which case the production specification is robust; a finding that the panic z-score profile flips with model capacity would localise the mechanism more narrowly to the chosen depth and learning rate.

Computational optimisation of stock selection The cross-sectional ranker used throughout this thesis (a 50-seed XGBoost ensemble at depth 4 with 500 trees per seed)

is treated as the canonical stock-picker for the analyses in Section 5.2, but it was not chosen by global optimisation. The coarse test-period sweep of Appendix F.3 and the architecture comparison sketched above leave open the likelihood that a more disciplined optimisation (finer hyperparameter grids under Bayesian search, end-to-end training against a portfolio-level objective rather than per-stock predicted return, or domain-specific architectures such as permutation-invariant set encoders that operate directly on the cross-section) would produce a strictly better-performing ranker. The qualitative findings of Section 5.2 are nevertheless robust to this concern: the regime-conditional shift in selection (continuation in the above-average tier, reversal in the below-average tier), the four-cluster heterogeneity in pick selection documented in Section 5.2.2, and the leg-beta amplification mechanism reflect properties of the data and the regime signal rather than artefacts of XGBoost’s specific optimum. A more optimal ranker would almost certainly improve absolute Sharpe, but the regime-conditioning story should survive that improvement. Quantifying both the absolute-performance gap left by the current specification and the stability of the qualitative findings across better-tuned rankers would tighten the link between the methodological choice and the strategic claim.

Exposure-scaling overlay The strategy’s primary vulnerability, discussed in Section 6.2, is a prolonged economic deterioration where beaten-down stocks do not recover and the reversal mechanism fails. XGB’s stock-selection channel does not address this directly: its contribution is to choose which stocks to hold, not how much. A natural extension is to combine it with a Daniel and Moskowitz (2016)-style exposure-scaling overlay that reduces gross exposure when a sustained-stress indicator is triggered, for example when π_t^{filter} has been above a threshold for a minimum number of consecutive months, or when its rolling maximum over a fixed window exceeds a calibrated bound. Preliminary checks on the 2011–2025 sample show no improvement from such an overlay because the sample contains no prolonged recession, so the circuit breaker only cuts profitable panic months; its value would emerge in a sustained downturn outside the current window. The dot-com sub-period of the 30-year expanding-window backtest (Section 5.3.3) is a natural setting to evaluate candidate overlay rules.

References

- James H. Albert and Siddhartha Chib. Bayes inference via Gibbs sampling of autoregressive time series subject to Markov mean and variance shifts. *Journal of Business & Economic Statistics*, 11(1):1–15, 1993. doi: 10.1080/07350015.1993.10509929.
- Andrew Ang and Geert Bekaert. International asset allocation with regime shifts. *The Review of Financial Studies*, 15(4):1137–1187, 2002. doi: 10.1093/rfs/15.4.1137.
- Clifford S. Asness, Tobias J. Moskowitz, and Lasse Heje Pedersen. Value and momentum everywhere. *The Journal of Finance*, 68(3):929–985, 2013. doi: 10.1111/jofi.12021.
- Nicholas Barberis, Andrei Shleifer, and Robert Vishny. A model of investor sentiment. *Journal of Financial Economics*, 49(3):307–343, 1998. doi: 10.1016/S0304-405X(98)00027-0.
- Pedro Barroso and Pedro Santa-Clara. Momentum has its moments. *Journal of Financial Economics*, 116(1):111–120, 2015. doi: 10.1016/j.jfineco.2014.11.010.
- Leonard E. Baum, Ted Petrie, George Soules, and Norman Weiss. A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains. *The Annals of Mathematical Statistics*, 41(1):164–171, 1970. doi: 10.1214/aoms/1177697196.
- Heiner Beckmeyer and Timo Wiedemann. All days are not created equal: Understanding momentum by learning to weight past returns. *Journal of Banking and Finance*, 181, 2025. doi: 10.1016/j.jbankfin.2025.107510.
- Daniele Bianchi, Andrea De Polis, and Ivan Petrella. Time-varying skewness and momentum crashes. *SSRN Working Paper*, 2022. doi: 10.2139/ssrn.4182040.
- Markus K. Brunnermeier and Lasse Heje Pedersen. Market liquidity and funding liquidity. *Review of Financial Studies*, 22(6):2201–2238, 2009. doi: 10.1093/rfs/hhn098.
- Tianqi Chen and Carlos Guestrin. XGBoost: A scalable tree boosting system. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 785–794, 2016. doi: 10.1145/2939672.2939785.
- Michael J. Cooper, Roberto C. Gutierrez Jr., and Allaudeen Hameed. Market states and momentum. *The Journal of Finance*, 59(3):1345–1365, 2004. doi: 10.1111/j.1540-6261.2004.00665.x.
- Kent Daniel and Tobias J. Moskowitz. Momentum crashes. *Journal of Financial Economics*, 122(2):221–247, 2016. doi: 10.1016/j.jfineco.2015.12.002.

- Werner F. M. De Bondt and Richard Thaler. Does the stock market overreact? *The Journal of Finance*, 40(3):793–805, 1985. doi: 10.1111/j.1540-6261.1985.tb05004.x.
- Arthur P. Dempster, Nan M. Laird, and Donald B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1):1–38, 1977. doi: 10.1111/j.2517-6161.1977.tb01600.x.
- Eugene F. Fama and Kenneth R. French. A five-factor asset pricing model. *Journal of Financial Economics*, 116(1):1–22, 2015. doi: 10.1016/j.jfineco.2014.10.010.
- Sylvia Frühwirth-Schnatter. *Finite Mixture and Markov Switching Models*. Springer, 2006. doi: 10.1007/978-0-387-35768-3.
- Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. *Bayesian Data Analysis*. Chapman and Hall/CRC, 3rd edition, 2013. doi: 10.1201/b16018.
- John Geweke. Bayesian treatment of the independent student- t linear model. *Journal of Applied Econometrics*, 8:S19–S40, 1993. doi: 10.1002/jae.3950080504.
- Christian L. Goulding, Campbell R. Harvey, and Michele G. Mazzoleni. Momentum turning points. *Journal of Financial Economics*, 149(3):378–406, 2023. doi: 10.1016/j.jfineco.2023.05.011.
- Shihao Gu, Bryan Kelly, and Dacheng Xiu. Empirical asset pricing via machine learning. *The Review of Financial Studies*, 33(5):2223–2273, 2020. doi: 10.1093/rfs/hhaa009.
- Massimo Guidolin and Allan Timmermann. Asset allocation under multivariate regime switching. *Journal of Economic Dynamics and Control*, 31(11):3503–3544, 2007. doi: 10.1016/j.jedc.2006.12.004.
- James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989. doi: 10.2307/1912559.
- Peter D. Hoff. *A First Course in Bayesian Statistical Methods*. Springer, 2009. doi: 10.1007/978-0-387-92407-6.
- Harrison Hong and Jeremy C. Stein. A unified theory of underreaction, momentum, and overreaction in asset markets. *The Journal of Finance*, 54(6):2143–2184, 1999. doi: 10.1111/0022-1082.00184.
- Narasimhan Jegadeesh. Evidence of predictable behavior of security returns. *The Journal of Finance*, 45(3):881–898, 1990. doi: 10.1111/j.1540-6261.1990.tb05110.x.

- Narasimhan Jegadeesh and Sheridan Titman. Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1):65–91, 1993. doi: 10.1111/j.1540-6261.1993.tb04702.x.
- Narasimhan Jegadeesh and Sheridan Titman. Profitability of momentum strategies: An evaluation of alternative explanations. *The Journal of Finance*, 56(2):699–720, 2001. doi: 10.1111/0022-1082.00342.
- Charles M.C. Lee and Bhaskaran Swaminathan. Price momentum and trading volume. *The Journal of Finance*, 55(5):2017–2069, 2000. doi: 10.1111/0022-1082.00280.
- Bruce N. Lehmann. Fads, martingales, and market efficiency. *The Quarterly Journal of Economics*, 105(1):1–28, 1990. doi: 10.2307/2937816.
- Scott M. Lundberg and Su-In Lee. A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc., 2017.
- Scott M. Lundberg, Gabriel Erion, Hugh Chen, Alex DeGrave, Jordan M. Prutkin, Bala Nair, Ronit Katz, Jonathan Himmelfarb, Nisha Bansal, and Su-In Lee. From local explanations to global understanding with explainable AI for trees. *Nature Machine Intelligence*, 2(1):56–67, 2020. doi: 10.1038/s42256-019-0138-9.
- Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952. doi: 10.1111/j.1540-6261.1952.tb01525.x.
- Stefan Nagel. Evaporating liquidity. *The Review of Financial Studies*, 25(7):2005–2039, 2012. doi: 10.1093/rfs/hhs066.
- Robert Novy-Marx. Is momentum really momentum? *Journal of Financial Economics*, 103(3):429–453, 2012. doi: 10.1016/j.jfineco.2011.05.003.
- Yizhan Shu and John M. Mulvey. Dynamic factor allocation leveraging regime-switching signals. *arXiv preprint arXiv:2410.14841*, 2024.
- Tyler Shumway. Forecasting bankruptcy more accurately: A simple hazard model. *The Journal of Business*, 74(1):101–124, 2001. doi: 10.1086/209665.
- Dimitri Vayanos and Paul Woolley. An institutional theory of momentum and reversal. *Review of Financial Studies*, 26(5):1087–1145, 2013. doi: 10.1093/rfs/hht014.

Appendix A: Data Sources and Variable Definitions

Three primary data sources are used. **CRSP** (Center for Research in Security Prices, via WRDS) provides monthly stock-level returns, prices, shares outstanding, exchange codes, and share codes from the Monthly Stock File (`crsp.msf`), as well as daily market returns from the Daily Stock Index file (`crsp.dsi`). **Compustat** (via WRDS) provides quarterly fundamental accounting data from `comp.fundq`, linked to CRSP via the CCM bridge table using primary link types (LU, LC). **FRED** (Federal Reserve Economic Data) provides Moody’s BAA and AAA corporate bond yields for computing the credit spread.

Following [Shumway \(2001\)](#), delisting returns are incorporated to prevent survivorship bias. When a stock is delisted, its actual delisting return is used if available; for performance-related delistings (CRSP codes 500–584) with missing returns, -30% is imputed.

A comprehensive overview of all data variables is provided in [Table 14](#).

Variable	Source	Description	Used in
<i>Stock-level variables</i>			
Monthly returns	CRSP <code>msf</code>	Adjusted holding-period returns	Sections 3, 5
Price, shares out.	CRSP <code>msf</code>	For market equity computation	Section 4
Exchange, share codes	CRSP <code>msf</code>	Universe filtering (codes 10, 11)	Section 3.2
Delisting returns	CRSP <code>msdelist</code>	Survivorship bias correction	Section 3.2
Book equity (<code>ceqq</code>)	Compustat <code>fundq</code>	Book-to-market, ROE	Appendix K
Total assets (<code>atq</code>)	Compustat <code>fundq</code>	Asset growth, gross profit, CF ratios	Appendix K
Income (<code>ibq</code>)	Compustat <code>fundq</code>	ROE, earnings growth, accruals	Appendix K
Revenue (<code>revtq</code>)	Compustat <code>fundq</code>	Gross profitability	Appendix K
COGS (<code>cogsq</code>)	Compustat <code>fundq</code>	Gross profitability	Appendix K
Debt (<code>dlttq</code> , <code>dlcq</code>)	Compustat <code>fundq</code>	Leverage ratio	Appendix K
Operating cash flow (<code>oancfy</code>)	Compustat <code>fundq</code>	CFO, FCF, accruals	Appendix K
Capital expenditure (<code>capxy</code>)	Compustat <code>fundq</code>	Free cash flow	Appendix K
<i>Market-level variables (HMM inputs)</i>			
Daily market returns	CRSP <code>dsi</code>	Drawdown (DD) computation	Section 4.1
Market price index	CRSP <code>dsi</code>	Drawdown (DD)	Section 4.1
Cross-sect. stock returns	CRSP <code>msf</code>	Dispersion (DISP), REL_N	Section 4.1
BAA, AAA bond yields	FRED	Credit spread (CS)	Section 4.1
<i>Derived features</i>			
$mom_1, \dots, mom_{12}$	Computed	12 trailing momentum signals	Section 4.4
π_t^{filter}	HMM output	Filtered panic probability	Section 4.3
10 fundamentals (ablation only)	Computed	bm, roe, gross profit, cash flow, etc.	Appendix K

Table 14: Summary of all data variables used in the analysis.

Notes: All data accessed via WRDS. Compustat linked to CRSP via the CCM bridge table using primary link types (LU, LC). Fundamentals are lagged four months from fiscal quarter end to ensure real-time availability. The sample covers January 1990 to November 2025 (training: 1990–2010; test: 2011–2025).

Appendix B: Portfolio Formation, Transaction Costs, and Performance Metrics

B.1 Portfolio Construction

For each method, a long-short portfolio is constructed each month. Stocks are ranked by their scores, and NYSE-listed stocks are used to determine decile breakpoints. All stocks (including AMEX and NASDAQ) whose score exceeds the 90th percentile breakpoint enter the long leg; all stocks below the 10th percentile breakpoint enter the short leg. Long-short construction is the standard in the momentum literature (Jegadeesh and Titman, 1993; Daniel and Moskowitz, 2016; Barroso and Santa-Clara, 2015) because it isolates momentum’s cross-sectional stock selection from market beta, providing a cleaner test of the strategy’s ability to distinguish future winners from losers.

Portfolio weights within each leg are assigned proportionally to market equity (value-weighting):

$$w_t^{s,L} = \frac{ME_t^s}{\sum_{s \in \mathcal{L}_t} ME_t^s}, \quad w_t^{s,S} = \frac{ME_t^s}{\sum_{s \in \mathcal{S}_t} ME_t^s}, \quad (29)$$

where $\mathcal{L}_t$ and $\mathcal{S}_t$ are the sets of stocks in the long and short legs at month t , respectively. The long-short portfolio return is:

$$r_t^{LS} = \sum_{s \in \mathcal{L}_t} w_t^{s,L} \cdot r_t^s - \sum_{s \in \mathcal{S}_t} w_t^{s,S} \cdot r_t^s. \quad (30)$$

Value-weighting tilts each leg toward large, liquid stocks, reducing the influence of small-cap return anomalies and improving implementability.

B.2 Transaction Costs

All portfolios are rebalanced monthly.

A one-way transaction cost of $c = 10$ basis points is applied to the turnover in both the long and short legs at each rebalancing date. The net portfolio return at month t is:

$$r_t^{\text{net}} = r_t^{\text{gross}} - c \cdot (\tau_t^L + \tau_t^S), \quad (31)$$

where τ_t^L and τ_t^S are the one-way turnover in the long and short legs, respectively:

$$\tau_t^L = \frac{1}{2} \sum_s \left| w_t^{s,L,\text{new}} - w_{t-1}^{s,L,\text{prev}} \right|, \quad \tau_t^S = \frac{1}{2} \sum_s \left| w_t^{s,S,\text{new}} - w_{t-1}^{s,S,\text{prev}} \right|, \quad (32)$$

and each summation runs over all stocks appearing in either the new or previous portfolio for that leg. The 10 basis point cost reflects the transaction environment for a value-

weighted, large-cap-biased strategy where NYSE breakpoints tilt toward liquid securities.

It is important to note that transaction costs are applied *post-hoc* as an accounting adjustment and are not incorporated into the models' training objectives. The cross-sectional models (DET, LR, and XGB) optimise stock-level prediction (classification accuracy or return forecasting) without knowledge of the previous portfolio, the resulting turnover, or the cost of rebalancing. This is standard practice in the empirical asset pricing literature (e.g., [Fama and French, 2015](#); [Goulding et al., 2023](#)), where portfolio construction is treated as a fixed downstream step applied uniformly to all strategies. An end-to-end framework that jointly optimises stock selection and portfolio-level objectives (e.g., net-of-cost Sharpe ratio) would require a fundamentally different architecture, such as a reinforcement learning or differentiable portfolio optimisation layer, and remains a potential extension of this work.

B.3 Performance Metrics

Standard Metrics Strategy performance is assessed using four standard risk-adjusted metrics computed on monthly net-of-fee returns.

Annualized return The geometric annualized return over n months is:

$$R_{\text{ann}} = \left[\prod_{t=1}^n (1 + r_t) \right]^{12/n} - 1. \quad (33)$$

Annualized volatility The annualized standard deviation of monthly returns:

$$\sigma_{\text{ann}} = \sigma(r) \cdot \sqrt{12}. \quad (34)$$

Sharpe ratio The annualized Sharpe ratio:¹⁵

$$SR = \frac{\bar{r}}{\sigma(r)} \cdot \sqrt{12}, \quad (35)$$

where $\bar{r}$ is the mean monthly return. Sharpe is computed from the arithmetic monthly mean annualised by $\sqrt{12}$, while annualised returns reported elsewhere in this thesis are compound (geometric); under high volatility the two definitions of average return differ, so Sharpe should not be inferred from a table by dividing annualised return by annualised volatility (e.g. at depth 1 in the methodology hyperparameter sweep, the reported Sharpe is 0.53 but $8.7\%/19.4\% = 0.45$).

¹⁵All Sharpe ratios in this thesis are computed using raw returns rather than excess returns above the risk-free rate. Since the same convention is applied to every strategy, relative comparisons and rankings are unaffected. The risk-free rate was near zero for most of the test period (2011–2021) and averaged roughly 2% annualized over the full 2011–2025 window, so absolute Sharpe ratios are overstated by approximately 0.05–0.10.

Maximum drawdown Let $C_t = \prod_{\tau=1}^t (1 + r_\tau)$ denote the cumulative wealth path. The maximum drawdown is:

$$MDD = \min_{1 \leq t \leq n} \frac{C_t - \max_{1 \leq \tau \leq t} C_\tau}{\max_{1 \leq \tau \leq t} C_\tau}. \quad (36)$$

Regime-Conditional Evaluation To assess whether the regime signal adds value beyond unconditional stock selection, performance is decomposed by market regime. Months are partitioned into calm ($\pi_t^{\text{filter}} < 0.5$) and panic ($\pi_t^{\text{filter}} \geq 0.5$) periods, and the Sharpe ratio is computed separately within each subset:

$$SR_{\text{calm}} = \frac{\bar{r}_{\text{calm}}}{\sigma(r_{\text{calm}})} \cdot \sqrt{12}, \quad SR_{\text{panic}} = \frac{\bar{r}_{\text{panic}}}{\sigma(r_{\text{panic}})} \cdot \sqrt{12}. \quad (37)$$

This decomposition helps assess whether the strategy’s alpha originates primarily from calm-period stock selection, panic-period risk avoidance, or both.

Appendix C: Feature Importance Methods

C.1 SHAP Values for XGBoost

To interpret the XGBoost model and assess the marginal contribution of each feature, SHAP (SHapley Additive exPlanations; [Lundberg and Lee, 2017](#)) values are computed on the test set. The SHAP value of feature j for observation $\mathbf{x}$ measures its average marginal contribution to the prediction across all possible feature orderings:

$$\phi_j(\mathbf{x}) = \sum_{S \subseteq N \setminus \{j\}} \frac{|S|! (p - |S| - 1)!}{p!} [f_{S \cup \{j\}}(\mathbf{x}) - f_S(\mathbf{x})], \quad (38)$$

where N is the full feature set, $p = |N|$, and $f_S(\mathbf{x})$ is the expected model output when only features in S are observed. For tree-based models, the TreeSHAP algorithm ([Lundberg et al., 2020](#)) computes these values exactly in polynomial time. The overall importance of each feature is summarised as:

$$I_j = \frac{1}{n_{\text{test}}} \sum_{i=1}^{n_{\text{test}}} |\phi_j(\mathbf{x}_i)|. \quad (39)$$

Two additional analyses are performed. First, the mean absolute SHAP value of π_t^{filter} is computed separately for calm and panic months to assess whether the regime signal’s influence varies across market states. Second, regime-specific XGBoost models are trained on calm-only and panic-only subsets of the training data, and their SHAP feature importances are compared to identify which features the model relies on differentially across regimes.

Appendix D: Alternative Training Targets for XGBoost

Section 5.3 argues that embedding a risk penalty directly in the XGBoost training target degrades performance, motivating the CRRA specification used in the main text. Table 15 shows this degradation is not an artefact of the mean-variance functional form: the same pattern appears under three families of training targets with varying risk tolerance.

Target family	Parameter	Sharpe	Ann. Ret	Ann. Vol	MDD
Baseline (r)	–	1.107	30.0%	27.1%	–28.3%
Mean-variance: $r - \frac{\gamma}{2}\sigma^2$	$\gamma = 0.5$	1.061	22.9%	21.6%	–22.3%
	$\gamma = 1$	1.024	16.5%	16.2%	–22.8%
	$\gamma = 2$	1.135	15.5%	13.6%	–22.0%
	$\gamma = 5$	1.062	13.5%	12.7%	–22.5%
	$\gamma = 10$	1.082	13.1%	12.1%	–18.9%
Log mean-variance: $\log(1+r) - \frac{\gamma}{2}\sigma^2$	$\gamma = 0.5$	1.051	14.6%	13.9%	–17.9%
	$\gamma = 1$	1.039	13.8%	13.4%	–23.8%
	$\gamma = 2$	1.080	14.0%	12.9%	–20.9%
	$\gamma = 5$	1.127	14.1%	12.4%	–18.7%
	$\gamma = 10$	0.995	12.1%	12.3%	–20.2%
Sharpe-like: r / σ^a	$a = 0.5$	1.148	24.8%	21.4%	–22.4%
	$a = 1.0$	0.984	14.6%	15.0%	–29.1%
	$a = 1.5$	0.936	12.4%	13.5%	–21.0%
	$a = 2.0$	0.759	12.8%	18.1%	–24.5%

Table 15: Out-of-sample performance of XGBoost (XGB) under alternative training targets (long-only construction). Every risk-adjusted target reduces returns faster than volatility.

The pattern is consistent: returns fall faster than volatility under every risk penalty, regardless of functional form. The mildest adjustments (mean-variance with $\gamma = 0.5$, Sharpe-like with $a = 0.5$) come closest to the baseline but still underperform. This confirms that the diagnostic in Section 5.3 is not specific to the mean-variance utility formulation; the single-model training-target approach is structurally unsuited to incorporating risk aversion in a long-short setting, which the CRRA dual-utility framework resolves.

Appendix E: CRRA Dual-Utility Construction

This appendix provides the technical specification behind the risk-aversion sweep in Section 5.3. The strategy uses two forecasts for each stock-month: an expected-return forecast $\hat{r}_t^s$ and a forward-volatility forecast $\hat{\sigma}_t^s$. $\hat{r}_t^s$ is produced by the XGBoost ensemble described in Section 3. $\hat{\sigma}_t^s$ is produced by a second XGBoost ensemble trained on the

realised forward volatility

$$\sigma_{[t+1, t+12]}^s = \text{std}(r_{t+1}^s, \dots, r_{t+12}^s),$$

using the same feature set augmented with trailing 12-month realised volatility.

For each stock, utility is computed under both trading directions:

$$\begin{aligned} u_{\text{long}}^s &= +\text{sign}(\hat{r}^s) \log(1 + |\hat{r}^s|/\varepsilon) - \gamma \log(\hat{\sigma}^s), \\ u_{\text{short}}^s &= -\text{sign}(\hat{r}^s) \log(1 + |\hat{r}^s|/\varepsilon) - \gamma \log(\hat{\sigma}^s), \end{aligned} \tag{40}$$

with $\varepsilon = 0.01$. The return term is $\pm \log$ -compressed expected return under each direction; the $-\gamma \log(\hat{\sigma})$ term enters with the same sign in both, so a low- σ stock is preferred whether considered as a long or as a short.

Each month, longs are the top decile of u_{long} and shorts are the top decile of u_{short} , both using NYSE breakpoints and full-universe membership. Positions are value-weighted within each leg. Stocks appearing in both top deciles (the overlap set, which grows with γ) are retained on both sides; their contributions to the long and short books net in $r_L - r_S$. This is the mechanism by which the long-short spread converges to zero as γ rises beyond the useful range: at high γ , the volatility penalty dominates the return signal, low- $\hat{\sigma}$ stocks occupy both top deciles simultaneously, and the two legs become near-identical value-weighted low-vol portfolios.

The 50-seed ensemble is trained on 1990–2010 and evaluated out-of-sample on 2011–2025. The sweep in Figure 11 uses a 17-point grid over $[0, 1]$ (step 0.05 up to 0.5, step 0.1 thereafter).

Appendix F: Robustness Checks

This section reports six robustness checks that complement the main results in Section 5.

F.1 Sub-Period Stability

To assess whether the main findings are driven by a single favourable sub-period, the test sample is split into three roughly equal windows: 2011–2015, 2016–2020, and 2021–2025.

	2011–2015	2016–2020	2021–2025	Full
Market	0.72	1.04	0.73	0.84
Fixed 12-mo	0.74	-0.20	-0.11	0.06
LR	0.47	-0.02	-0.38	-0.01
XGB	0.59	1.04	1.70	1.11

Table 16: Sub-period Sharpe ratios (long-short). The test period is split into three roughly equal sub-periods. XGB is the only strategy with positive Sharpe ratios in all three sub-periods.

In long-short construction, XGB is the only strategy that produces positive Sharpe ratios in every sub-period (0.59, 1.04, 1.70). LR collapses in 2016–2020 (Sharpe -0.02) and turns negative in 2021–2025 (-0.38). Fixed 12-month momentum turns negative after 2015, confirming the momentum crash phenomenon documented by [Daniel and Moskowitz \(2016\)](#). XGB’s strongest performance (1.70) is in 2021–2025, a period characterised by elevated volatility and multiple stress episodes, precisely the conditions where regime-conditional stock selection adds the most value.

F.2 Turnover and Transaction Cost Sensitivity

A strategy’s net-of-cost performance depends critically on its turnover. Table 17 reports average monthly one-way turnover for each strategy.

Strategy	Avg. Monthly TO	Annualised TO
Fixed 12-mo	70.1%	842%
Fixed 1-mo	184.3%	2,211%
LR	168.8%	2,026%
XGB	132.3%	1,588%

Table 17: Portfolio turnover by strategy (long-short). Average monthly one-way turnover and annualised turnover (monthly $\times 12$).

In long-short construction, XGB’s average monthly one-way turnover is 132.3%, reflecting the regime-conditional changes in stock selection that shift the portfolio composition substantially. Table 18 translates these turnover differences into Sharpe ratio sensitivity across a range of transaction cost assumptions.

	0 bps	5 bps	10 bps	20 bps	30 bps	50 bps
Fixed 12-mo	0.10	0.08	0.07	0.03	0.00	-0.06
Fixed 1-mo	0.39	0.32	0.26	0.14	0.01	-0.24
LR	0.12	0.05	-0.02	-0.16	-0.29	-0.57
XGB	1.19	1.15	1.11	1.03	0.95	0.78

Table 18: XGB remains profitable at 50 bps. All other strategies become unprofitable at moderate cost levels.

XGB remains profitable across all cost levels tested: even at 50bps one-way (five times the baseline), its Sharpe ratio is 0.78. All other strategies become unprofitable at moderate cost levels. The baseline 10bps assumption is reasonable for a value-weighted strategy using NYSE breakpoints, which tilts toward the most liquid names.

F.3 Model Sensitivity

The XGBoost model uses hyperparameters selected via cross-validation on the training set. To verify that the results are not an artefact of a particular configuration, Table 19 reports out-of-sample performance under alternative tree depths, learning rates, and ensemble sizes. A depth-matched Random Forest baseline is discussed below; full numbers are in `results/thesis/random_forest_results.csv`.

Configuration	Sharpe	Ann. Ret	Ann. Vol
depth=3, lr=0.05, $n=500$	0.98	19.2%	19.9%
depth=4, lr=0.05, $n=500$	1.11	21.7%	19.5%
depth=5, lr=0.05, $n=500$	0.92	17.4%	19.4%
depth=6, lr=0.05, $n=500$	0.90	17.0%	19.5%
depth=4, lr=0.01, $n=500$	0.99	18.7%	19.2%
depth=4, lr=0.10, $n=500$	1.00	19.6%	19.9%
depth=4, lr=0.05, $n=200$	1.06	20.5%	19.4%
depth=4, lr=0.05, $n=1000$	0.96	18.7%	19.9%

Table 19: XGBoost hyperparameter sensitivity (long-short). Each row varies one hyperparameter from the baseline. The baseline is not the single best configuration, but all variants outperform unconditional momentum.

Hyperparameter sensitivity. At the baseline depth of 4, the Sharpe ratio ranges from 0.89 to 1.11 across learning rate and ensemble size variations. Performance degrades at depths below 3 (see Section 3.2.1 for the depth analysis) and above 5 due to overfitting. The results confirm that XGB’s performance is not an artefact of a particular hyperparameter configuration.

Boosting versus bagging. As a complementary check, a Random Forest matched on depth (4), ensemble size (500 trees per seed), and feature set (the same 13 features) was trained over 50 seeds (`results/thesis/random_forest_results.csv`). Sharpe falls from 1.08 to 0.73 and π_t^{filter} ’s SHAP share drops from 45% to 18%. The mechanism is the difference between boosting and bagging. XGBoost fits each successive tree to the residuals of the current ensemble, so at every iteration the algorithm preferentially selects whichever feature most reduces the current error; for this problem that feature is π_t^{filter} , and its importance compounds across trees. Random Forest grows each tree independently on a bootstrap sample and subsamples a random subset of features at every split, which

by construction prevents any single feature from dominating the ensemble and distributes importance more evenly across the 13 features. Nonlinearity and depth are therefore necessary but not sufficient: the boosted additive structure is what converts the regime signal into a primary gating variable rather than a marginal predictor, which is the basis for the architecture-dependence claim in Section 3.2.1.

F.4 Number of Hidden States

The baseline HMM uses $K = 2$ states (calm and panic). A natural question is whether finer regime granularity (for example, distinguishing mild corrections from severe crises) would improve the downstream portfolio models. Table 20 reports out-of-sample Sharpe ratios for $K = 2, 3, 4, 5$, each averaged over five representative Gibbs sampler seeds (the production regime signal averages 200 seeds, but five suffice for this diagnostic).

K	LR		XGB	
	Mean Sharpe	Std	Mean Sharpe	Std
2	-0.012	0.000	1.107	0.000
3	-0.029	0.004	0.831	0.164
4	-0.029	0.003	0.606	0.291
5	-0.028	0.003	0.724	0.087

Table 20: Out-of-sample Sharpe ratios for HMMs with $K = 2, 3, 4, 5$ states (long-short). Each entry reports the mean and standard deviation across 5 independent Gibbs sampler seeds (2,000 iterations each). LR is invariant to K . XGB shows no consistent improvement beyond $K = 2$, and cross-seed variability increases with K .

In long-short construction, XGB shows no consistent improvement beyond $K = 2$. Cross-seed variability increases with K , and the intermediate states that emerge for $K \geq 3$ lack the persistence needed to provide a stable conditioning signal. With only 241 training months, the data cannot reliably separate more than two regimes for the purpose of downstream cross-sectional stock selection.

Beyond estimation, $K = 2$ has an interpretability advantage. The two-state model produces a single continuous signal (π_t^{filter}) that XGBoost can condition on in a straightforward way: higher values indicate more stress. With $K \geq 3$, the additional states may capture dimensions of the economy that are orthogonal to the continuation-reversal mechanism, for example a “growth” state characterised by strong fundamentals but moderate stress indicators. Such a state would conflate two economically distinct situations (a genuine expansion vs. a recovery from panic) and could introduce noise into the regime signal, degrading the model’s ability to distinguish continuation from reversal. The two-state framework avoids this by restricting the regime classification to a single dimension, the degree of market stress, which is the dimension most relevant to momentum’s cross-sectional structure.

F.5 January Exclusion

Momentum is known to reverse in January (Jegadeesh and Titman, 1993): past losers outperform past winners, partially offsetting momentum’s annual profitability. To verify that XGB’s performance is not driven by a January seasonal, Table 21 recomputes all metrics after excluding the 14 January months from the test period. XGB’s Sharpe ratio declines marginally from 1.11 to 1.06, confirming that the regime-momentum interaction operates throughout the calendar year. Fixed 12-month momentum improves slightly when January is excluded (0.06 to 0.14), consistent with the well-documented January reversal in unconditional momentum.

	Ann. Ret	Ann. Vol	Sharpe	Months
<i>Full sample (baseline)</i>				
XGB	21.7%	19.5%	1.11	167
Fixed 12-mo mom	−1.9%	26.1%	0.06	167
<i>Excluding January</i>				
XGB	20.9%	19.9%	1.06	153
Fixed 12-mo mom	0.1%	25.6%	0.14	153

Table 21: January exclusion test. Performance is recomputed after dropping all 14 January months from the test period.

Appendix G: External Validity

G.1 Alternative Train/Test Splits

Table 22 reports the baseline train/test split performance. The 30-year expanding-window backtest (Section 5.3.3) provides further out-of-sample evidence across alternative training windows and prediction horizons.

Train / Test split	N_{test}	LR Sharpe	XGB Sharpe
1990–2010 / 2011–2025 (baseline)	167	−0.01	1.11

Table 22: Out-of-sample performance under the baseline train/test split (long-short, mom+ π , 50-seed ensemble). LR collapses while XGB achieves a Sharpe of 1.11 with highly significant factor model alphas.

In the baseline split (1990–2010 training, 2011–2025 test), XGB achieves a Sharpe of 1.11 while LR collapses (−0.01), confirming the nonlinear regime-momentum interaction documented in Section 5.1.

G.2 30-Year Expanding-Window Out-of-Sample Methodology

The historical evidence reported in Section 5.3.3 comes from a separate analysis from the one above: a 30-year out-of-sample backtest in which the full pipeline (HMM + XGBoost) is re-estimated every January from 1995 onward, with each year’s predictions made by a model trained only on data available through the prior December. Each retraining uses the production seed configuration: 200 HMM seeds (Gibbs sampler with 2000 iterations, 500 burn-in), 50 XGBoost seeds (depth 4, learning rate 0.05, 500 trees), and the same long-short construction described in Appendix B.1.

The HMM’s panic-state label is fixed by an economically motivated sign convention rather than the data-driven percentile comparison used elsewhere. The convention encodes the known direction of each feature in panic regimes: -1 for drawdown, $+1$ for cross-sectional dispersion, -1 for relative market participation, and $+1$ for the credit spread. The data-driven label-switching procedure used in the main pipeline is unstable on short training windows with few observed crisis months (under 20), so the fixed-sign fallback activates for retrains before 2002. From 2002 onward the data-driven procedure resumes, by which point the training window contains the dot-com bust as a calibration anchor. The HMM seed loop is parallelised across CPU cores via `multiprocessing.Pool`; the parallel implementation is bit-identical to the serial reference (verified in the test suite by comparing per-seed pi-filter arrays).

Each retraining produces 12 monthly out-of-sample long-short returns, which are stitched into a single 30-year monthly time series spanning January 1995 to November 2024 (359 months). Every prediction is strictly causal: the model used for month t saw no information after month $t - 1$. The 167-month sub-window from January 2011 onward overlaps the main test period reported in Section 5.1 and serves as a methodological cross-check; expanding-window OOS produces a Sharpe of 0.88 over that window, within the production result’s bootstrap interval.

G.3 Factor Model Regressions

To test whether XGB’s out-of-sample return is explained by known risk factors, monthly returns are regressed on five standard factor models: CAPM, Fama–French three-factor, Carhart four-factor, Fama–French five-factor, and a six-factor specification that augments FF5 with the momentum factor (UMD). Newey–West standard errors (6 lags) are used throughout. Table 23 reports the intercept (alpha) and its t -statistic under each specification; Table 24 repeats the exercise for the fund-augmented variant. Alphas that survive all five factor models indicate return that cannot be attributed to market, size, value, profitability, investment, or momentum exposure.

Model	α (%)	$t(\alpha)$	Mkt-RF	SMB	HML	UMD	RMW	CMA
CAPM	23.4	4.08	-0.14					
FF3	23.2	4.48	-0.14	+0.04	-0.17			
Carhart	24.3	4.75	-0.19	-0.01	-0.22	-0.20		
FF5	22.9	4.62	-0.13	+0.08	-0.23		+0.06	+0.12
FF6	24.1	4.81	-0.18	+0.01	-0.31	-0.22	+0.01	+0.21

Table 23: Factor model regressions for XGB (mom+ π). α is annualised. t -statistics use Newey–West standard errors (6 lags).

For completeness, Table 24 reports the same regressions for the fund-augmented variant (XGB: mom+ π +fund) discussed in Section 5.1.3. Alphas remain statistically significant under all five specifications ($t > 2.9$) but are materially smaller than the baseline’s, consistent with the dilution documented in the main text.

Model	α (%)	$t(\alpha)$	Mkt-RF	SMB	HML	UMD	RMW	CMA
CAPM	17.6	3.04	-0.16					
FF3	17.1	3.12	-0.14	-0.06	-0.13			
Carhart	18.2	3.22	-0.19	-0.12	-0.18	-0.20		
FF5	16.4	3.02	-0.11	+0.01	-0.30		+0.09	+0.35
FF6	17.7	3.17	-0.16	-0.06	-0.39	-0.23	+0.04	+0.44

Table 24: Factor model regressions for XGB (mom+ π +fund). α is annualised. Newey–West t -statistics (6 lags).

G.4 Block Bootstrap Inference

To assess the sampling variability of XGB’s Sharpe ratio and test its outperformance formally, a block bootstrap is applied to the monthly return series. Returns are resampled with replacement in 12-month blocks to preserve autocorrelation, producing 10,000 synthetic 167-month paths (fixed seed 42). The Sharpe ratio is recomputed on each resample to obtain a 95% confidence interval, and the same block draws are used for XGB and each benchmark to form a paired test of the Sharpe difference. The bottom panel repeats the exercise on XGB returns split by regime.

	Sharpe	95% CI	vs XGB (p)
XGB	1.11	[0.67, 1.53]	—
LR	−0.01	[−0.41, 0.42]	0.001***
DET (Formula)	0.11	[−0.36, 0.59]	0.008***
Fixed 12-mo mom	0.06	[−0.43, 0.59]	0.005***
Fixed 1-mo mom	0.26	[−0.21, 0.73]	0.021**
<i>XGB regime-conditional Sharpe</i>			
XGB Calm ($n = 108$)	0.84	[0.29, 1.32]	—
XGB Panic ($n = 59$)	1.54	[0.85, 2.22]	—
Panic − Calm	0.70	[−0.16, 1.59]	0.109

Table 25: Block bootstrap confidence intervals for Sharpe ratios and paired tests of XGB’s outperformance. Uses 12-month blocks and 10,000 resamples with a fixed seed. The “vs XGB (p)” column reports two-sided p -values from paired tests where both strategies are resampled on the same dates. * $p < 0.10$; ** $p < 0.05$; *** $p < 0.01$.

XGB’s Sharpe confidence interval is [0.67, 1.53], comfortably above zero, while every benchmark’s interval straddles zero. The paired tests reject equality of Sharpe ratios between XGB and each benchmark at $p < 0.025$. Within XGB, the panic-period Sharpe exceeds the calm-period Sharpe in point estimate (1.54 vs 0.84), but with only 59 panic months the paired difference is not statistically significant ($p = 0.109$); the 95% interval [−0.16, 1.59] straddles zero. The point-estimate gap is consistent with the regime-conditional selection mechanism described in Section 5.2, but its statistical confirmation would require a longer panic sample. These intervals describe sampling variability within the 2011–2025 sample; out-of-sample generalisation is addressed separately by the 30-year expanding-window backtest (Section 5.3.3) and the sub-period analysis (Appendix F.1).

G.5 Leg-Level CAPM Betas by Regime

The D&M comparison in Section 5.3.2 characterises XGB’s contribution as stock selection rather than exposure scaling. A direct test of this framing is to compute leg-level CAPM betas separately by regime: if XGB reorganises which stocks it holds across regimes (cross-sectional re-ranking), the long-leg and short-leg betas should shift between calm and panic in opposite directions; if instead XGB mechanically scales exposure, both legs’ betas should move together.

Table 26: Leg-level CAPM betas by regime, 2011–2025 OOS test period. Panic months are those with $\pi_t^{\text{filter}} \geq 0.5$. Newey–West standard errors (3 lags) in parentheses.

Strategy	Leg	Full	Calm	Panic
Fixed 12-mo mom	Long	1.09 (0.06)	1.09 (0.09)	1.07 (0.08)
Fixed 12-mo mom	Short	1.74 (0.11)	1.80 (0.21)	1.69 (0.13)
XGB	Long	1.55 (0.08)	1.40 (0.11)	1.65 (0.10)
XGB	Short	1.10 (0.06)	1.22 (0.09)	0.98 (0.05)

The pattern in Table 26 is consistent with the cross-sectional re-ranking interpretation. XGB’s long-leg beta rises from 1.40 in calm to 1.65 in panic (*tilt amplification*: in panic, the long leg holds names with higher market sensitivity), while the short-leg beta falls from 1.22 to 0.98. Fixed 12-month momentum, by contrast, holds both leg-level betas roughly stable across regimes (long: 1.09 calm, 1.07 panic; short: 1.80 calm, 1.69 panic), consistent with the unconditional rank-based construction. The difference is direct evidence that the regime signal acts on *which* stocks are held rather than on *how much*.

G.6 Per-cluster Feature Signature ($K = 4$ Partition)

Section 5.2.2 reports a $K = 4$ partition of the 167 test months and characterises each cluster narratively. Table 27 provides the full per-cluster feature signature for completeness: cross-section context features (regime composition, momentum landscape, dispersion, skewness, recent strategy performance) computed at the month level and averaged across each cluster’s months; long-leg pick raw momentum at each horizon, averaged across each cluster’s months; and long-leg pick cross-sectional z-score (deviation from the cross-section mean at each horizon, in units of cross-section std) averaged across cluster months. Per-cluster Sharpes and other strategy outcomes appear in Table 10 in the main text.

Feature	Cluster 1 <i>calm cont.</i>	Cluster 2 <i>mild cont.</i>	Cluster 3 <i>post-panic</i>	Cluster 4 <i>deep crisis</i>
<i>Cross-section context</i>				
Cross-section mom (overall mean)	13.8%	7.3%	13.5%	−9.3%
Cross-section mom (short tertile, $h = 1-4$)	6.1%	4.4%	3.0%	−7.1%
Cross-section mom (mid tertile, $h = 5-8$)	14.6%	7.6%	12.4%	−11.1%
Cross-section mom (long tertile, $h = 9-12$)	20.8%	9.7%	25.0%	−9.6%
Cross-section dispersion (short)	0.23	0.27	0.28	0.23
Cross-section dispersion (mid)	0.44	0.46	0.56	0.35
Cross-section dispersion (long)	0.60	0.59	0.85	0.45
Cross-section skewness (short)	4.6	7.2	6.4	3.7
Cross-section skewness (mid)	6.2	7.6	6.2	5.6
Cross-section skewness (long)	6.7	8.1	6.5	5.7
π^{filter} frequency, prior 6 months	0.27	0.37	0.30	0.51
π^{filter} frequency, prior 12 months	0.31	0.32	0.41	0.30
Strategy Sharpe, prior 12 months	1.22	1.04	1.15	1.08
<i>Long-leg pick raw momentum by horizon</i>				
Long-leg pick mom, $h = 1$	7.1%	3.1%	−3.4%	−15.3%
Long-leg pick mom, $h = 2$	13.5%	8.6%	−5.5%	−21.6%
Long-leg pick mom, $h = 3$	18.3%	8.8%	−5.3%	−26.3%
Long-leg pick mom, $h = 4$	23.1%	10.6%	−5.8%	−30.8%
Long-leg pick mom, $h = 5$	30.8%	13.8%	−5.9%	−35.0%
Long-leg pick mom, $h = 6$	40.2%	18.0%	−3.2%	−38.3%
Long-leg pick mom, $h = 7$	49.4%	20.7%	−1.3%	−42.1%
Long-leg pick mom, $h = 8$	59.2%	23.7%	0.0%	−44.9%
Long-leg pick mom, $h = 9$	61.3%	23.1%	1.1%	−46.5%
Long-leg pick mom, $h = 10$	60.7%	21.9%	1.3%	−47.2%
Long-leg pick mom, $h = 11$	61.1%	22.8%	1.2%	−47.3%
Long-leg pick mom, $h = 12$	57.9%	18.5%	−5.3%	−47.5%
<i>Long-leg pick cross-sectional z-score by horizon</i>				
Long-leg z-curve, $h = 1$	0.26	0.04	−0.28	−0.73
Long-leg z-curve, $h = 2$	0.38	0.14	−0.31	−0.74
Long-leg z-curve, $h = 3$	0.40	0.11	−0.34	−0.75
Long-leg z-curve, $h = 4$	0.44	0.10	−0.36	−0.80
Long-leg z-curve, $h = 5$	0.54	0.16	−0.35	−0.85
Long-leg z-curve, $h = 6$	0.65	0.23	−0.29	−0.83
Long-leg z-curve, $h = 7$	0.73	0.27	−0.26	−0.85
Long-leg z-curve, $h = 8$	0.80	0.32	−0.22	−0.87
Long-leg z-curve, $h = 9$	0.76	0.29	−0.23	−0.88
Long-leg z-curve, $h = 10$	0.72	0.24	−0.25	−0.87
Long-leg z-curve, $h = 11$	0.67	0.22	−0.26	−0.86
Long-leg z-curve, $h = 12$	0.54	0.12	−0.35	−0.87

Table 27: Per-cluster feature signature for the $K = 4$ partition, 2011–2024 test period. Features are grouped into three blocks: (1) cross-section context features (regime composition, momentum landscape, dispersion, skewness, recent strategy performance) computed at the month level and averaged across each cluster’s months; (2) long-leg pick raw momentum at each horizon, averaged across each cluster’s months; (3) long-leg pick cross-sectional z-score (deviation from the cross-section mean at each horizon, in units of cross-section std) averaged across cluster months. Per-cluster Sharpes and other strategy outcomes appear in Table 10; this table is the full descriptive feature signature.

G.7 Cluster-Level Cross-Validation

Three diagnostics extend the regime-aggregate findings of Section 5.2 across the $K = 4$ partition.

Leg-betas across the $K = 4$ partition XGB retains long-leg $\beta >$ short-leg β in every cluster (spreads +0.68, +0.40, +0.49, +0.43 for clusters 1–4), while unconditional 12-month momentum exhibits leg-beta inversion in every cluster (−0.24, −0.87, −0.75, −0.42; Table 28). Daniel and Moskowitz (2016)’s crash mechanism is latent across every cluster type; XGB escapes it everywhere through compositional re-ranking.

Cluster	n	XGB			Fixed 12-mo mom		
		β_L	β_S	Spread	β_L	β_S	Spread
1 (calm continuation)	53	1.65 (0.22)	0.97 (0.13)	0.68	1.32 (0.22)	1.57 (0.24)	−0.24
2 (mild continuation)	62	1.60 (0.15)	1.20 (0.09)	0.40	1.04 (0.08)	1.91 (0.17)	−0.87
3 (post-panic recovery)	31	1.46 (0.16)	0.97 (0.11)	0.49	1.09 (0.17)	1.84 (0.22)	−0.75
4 (deep crisis)	21	1.49 (0.08)	1.06 (0.09)	0.43	1.07 (0.10)	1.49 (0.14)	−0.42

Table 28: Leg-level CAPM betas by K=4 cluster, 2011–2024 test period. XGB (regime-conditional XGBoost ensemble) versus unconditional 12-month momentum, side-by-side. Standard errors in parentheses. Spread = $\beta_L - \beta_S$.

π_t^{filter} ’s SHAP share by cluster Combined long-and-short SHAP shares are 0.50, 0.46, 0.54, 0.24 for clusters 1–4 (Table 29). The pattern is consistent with the routing interpretation: π_t^{filter} contributes most where the cross-section is ambiguous (cluster 3) and least where it is unambiguously crisis-shaped (cluster 4, where momentum splits alone suffice to produce the loser portfolio).

Cluster	n	Long π share	Short π share	Combined π share
1 (calm continuation)	53	34.0%	50.1%	49.9%
2 (mild continuation)	62	26.9%	46.5%	46.0%
3 (post-panic recovery)	31	27.8%	53.7%	53.6%
4 (deep crisis)	21	11.4%	38.1%	24.4%

Table 29: π_t^{filter} ’s SHAP share by K=4 cluster, 2011–2024 test period. Long and short shares are computed over long-leg and short-leg picks respectively; combined is over all stock-months in the cluster.

Short-leg structure (future work) The $K = 4$ partition was defined on long-leg z-curves; applying the same labels to the short leg shows the short basket is structurally distinct from a negated long-leg, with short-vs-negated-long correlations of +0.94, +0.89, +0.19, and −0.84 for clusters 1–4 (Table 30). The month-1 short-leg spike in cluster 4 (+0.57) replicates the asymmetric short-side sensitivity documented by Nagel (2012),

consistent with Section 5.2.1: month 1 carries 10.6% of panic short-leg momentum SHAP versus 5.3% in calm. A full short-leg analysis is left to future work.

Cluster	n	z_1	$\bar{z}$	Basket	$\rho(-z^L)$
1 (calm continuation)	53	-0.062	-0.365	391	0.943
2 (mild continuation)	62	0.101	-0.174	334	0.888
3 (post-panic recovery)	31	0.209	0.241	382	0.190
4 (deep crisis)	21	0.568	0.170	353	-0.835

Table 30: Short-leg z-curve summary by K=4 cluster, 2011–2024 test period. z_1 is the cluster-mean cross-sectional z-score of short-leg picks at the 1-month horizon; $\bar{z}$ averages across all 12 horizons. Basket size is mean stocks per month. $\rho(-z^L)$ is the correlation between the short-leg z-curve centroid and the negated long-leg z-curve centroid; values below 0.95 indicate structural asymmetry.

G.8 Ridge Regression Baseline

To isolate the effect of linearity from the choice of training target, a Ridge regression is estimated on the same 13 standardised features as LR and XGB, predicting continuous forward returns (the same target as XGBoost). Table 31 reports results across a range of regularisation strengths.

Model	α	Sharpe	Ann. Ret	Ann. Vol	MDD	Calm / Panic SR
OLS (no reg.)	—	-0.58	-11.1%	17.4%	-82.0%	-0.61 / -0.56
Ridge	0.1	-0.58	-11.1%	17.4%	-82.0%	-0.61 / -0.56
Ridge	1	-0.58	-11.1%	17.4%	-82.0%	-0.61 / -0.56
Ridge	10	-0.58	-11.1%	17.4%	-82.0%	-0.61 / -0.56
Ridge	100	-0.58	-11.0%	17.4%	-81.9%	-0.60 / -0.56
Ridge	1,000	-0.54	-10.5%	17.5%	-80.3%	-0.60 / -0.47
LR (LR, classification)	—	-0.01	-1.3%	15.2%	-42.4%	-0.29 / 0.39
XGB	—	1.11	21.7%	19.5%	-22.8%	0.84 / 1.53

Table 31: Ridge regression baseline predicting continuous forward returns (same target as XGB). The Sharpe of -0.57 is robust across all regularisation strengths.

The Ridge regression assigns π_t^{filter} a coefficient of +0.003, which is positive but economically negligible: a one-standard-deviation increase in the panic probability shifts the predicted return by only 0.3 basis points. The linear model cannot use the regime signal to condition how the momentum horizons jointly determine the stock ranking, regardless of the training target.

G.9 Emission Distribution Robustness

To verify that regime identification is robust to the choice of emission distribution, the HMM is re-estimated with multivariate Student- t emissions across a grid of degrees of

freedom $\nu \in \{3, 5, 7, 10, 15, 20, 30, 50, 100\}$.

The Student- t model is estimated via data augmentation (Geweke, 1993): each observation receives a latent weight $w_t \sim \text{Gamma}(\nu/2, \nu/2)$, so that $\mathbf{z}_t \mid s_t=k, w_t \sim \mathcal{N}(\boldsymbol{\mu}_k, \boldsymbol{\Sigma}_k/w_t)$. Marginalising over w_t yields a multivariate t with ν degrees of freedom. Crucially, conditional on w_t the emission is Gaussian, so the Normal-Inverse-Wishart conjugate updates carry over with weighted sufficient statistics, preserving the efficiency of the Gibbs sampler.

ν	LL (train)	LL (test)	BIC	Agreement	Corr(π)
3	−882.7	−701.1	1929.9	91.4%	0.872
5	−829.0	−669.8	1822.5	96.8%	0.960
7	−838.3	−698.2	1841.2	98.0%	0.984
10	−815.8	−644.3	1796.1	96.8%	0.960
15	−800.1	−651.0	1764.8	98.3%	0.997
20	−802.9	−621.7	1770.4	99.3%	0.997
30	−802.6	−640.6	1769.8	98.5%	0.993
50	−796.6	−658.1	1757.7	98.8%	0.996
100	−792.3	−643.5	1749.0	98.8%	0.987
Normal	−787.5	−642.5	1739.5	—	—

Table 32: Student- t vs. Normal emission HMM comparison. LL = log-likelihood; Agreement = fraction of months with identical binary regime classification as the Normal HMM; Corr(π) = correlation of filtered panic probabilities.

The Normal model achieves the lowest BIC (2257.3 vs. 2257.6 for the best Student- t at $\nu = 100$), indicating no meaningful improvement from heavier tails. Regime classifications agree with the Normal HMM on at least 97.8% of months across all ν values, and the correlation of filtered panic probabilities exceeds 0.98 everywhere. These results confirm that the Gaussian emission specification is robust for this application.

Appendix H: Label Switching: Sign Correction Procedure

Because the four HMM features respond to stress in different directions, the expected direction of each feature during stress is determined empirically. For each feature j , the 95th percentile is computed separately on training months falling within known crisis windows (the dot-com bust of 2000–2002 and the Global Financial Crisis of 2007–2009) and on all remaining training months:

$$\text{sign}_j = \begin{cases} +1 & \text{if } q_{0.95}^{\text{crisis}}(z_j) \geq q_{0.95}^{\text{normal}}(z_j), \\ -1 & \text{otherwise.} \end{cases} \quad (41)$$

Features that spike upward during crises receive $\text{sign}_j = +1$; features that decline receive $\text{sign}_j = -1$. The panic state is then identified as the regime $k \in \{0, 1\}$ whose posterior mean vector $\hat{\boldsymbol{\mu}}_k$ best aligns with the crisis direction:

$$k^* = \arg \max_{k \in \{0, 1\}} \sum_{j=1}^D \text{sign}_j \cdot \hat{\mu}_{k,j}. \quad (42)$$

The state labelled k^* is designated as “panic” and the other as “calm.”

Appendix I: Convergence Diagnostics




Figure 12: Trace plots of the Gibbs sampler. Each panel shows sampled regime means (calm in blue, panic in red) for one HMM feature. The dashed line marks the end of burn-in.

I.1 Gelman–Rubin Diagnostic

The Gelman–Rubin $\hat{R}$ statistic (Gelman et al., 2013) assesses convergence by comparing within-chain and between-chain variance across the five independent Gibbs sampler chains. For each parameter θ , $\hat{R} = \sqrt{\hat{V}/W}$, where W is the mean within-chain variance

and $\hat{V} = \frac{n-1}{n}W + \frac{B}{n}$ combines within-chain variance with the between-chain variance B . Values near 1.0 indicate that the chains have converged to the same posterior distribution; $\hat{R} < 1.1$ is the standard threshold.

Parameter	$\hat{R}$	Status
$\mu_{\text{panic}}(\text{DD})$	1.000	OK
$\mu_{\text{panic}}(\text{DISP})$	1.000	OK
$\mu_{\text{panic}}(\text{REL_N})$	1.000	OK
$\mu_{\text{panic}}(\text{CS})$	1.000	OK
$\mu_{\text{calm}}(\text{DD})$	1.000	OK
$\mu_{\text{calm}}(\text{DISP})$	1.001	OK
$\mu_{\text{calm}}(\text{REL_N})$	1.000	OK
$\mu_{\text{calm}}(\text{CS})$	1.000	OK
P_{00} (calm persistence)	1.000	OK
P_{01} (calm $\rightarrow$ panic)	1.000	OK
P_{10} (panic $\rightarrow$ calm)	1.000	OK
P_{11} (panic persistence)	1.000	OK

Table 33: Gelman–Rubin $\hat{R}$ statistics for all HMM parameters, computed across 5 independent chains (1,500 post-burn-in draws each). All values are below 1.01, well within the convergence threshold of 1.1.

I.2 Seed Convergence of the Regime Signal

Gelman–Rubin and the trace plots assess whether a single Gibbs chain has converged to its stationary posterior, but they do not quantify how many chains are needed to stabilise the *downstream* Sharpe ratio. The regime signal π_t^{filter} enters the pipeline as the average of N_s seed-specific forward filters (Section 3.1.5), and because XGBoost is a nonlinear function of π_t^{filter} , MCMC noise in any individual chain does not cancel linearly.

To measure this dependence, we subsample $k \in \{1, 2, 3, 5, 10, 15, 20, 30, 50, 75, 100\}$ chains without replacement from the seed pool, average the forward filter over those k chains, retrain the XGBoost ensemble on the resulting signal, and record the out-of-sample Sharpe. For each $k < 100$ we repeat on 30 random subsets; $k = 100$ exhausts the pool.

k	n_{sub}	Mean	Median	Std	P25	P75	Min	Max
1	30	1.040	1.031	0.078	0.985	1.090	0.915	1.205
2	30	1.062	1.081	0.069	1.019	1.103	0.894	1.217
3	30	1.049	1.044	0.043	1.014	1.075	0.973	1.143
5	30	1.042	1.042	0.047	1.002	1.083	0.967	1.154
10	30	1.053	1.050	0.036	1.038	1.066	0.965	1.145
15	30	1.075	1.079	0.029	1.055	1.097	1.008	1.139
20	30	1.078	1.080	0.022	1.065	1.090	1.025	1.122
30	30	1.089	1.082	0.025	1.069	1.109	1.052	1.137
50	30	1.093	1.092	0.015	1.082	1.101	1.067	1.124
75	30	1.098	1.102	0.012	1.091	1.105	1.067	1.118
100	1	1.118	1.118	0.000	1.118	1.118	1.118	1.118

Table 34: XGBoost ensemble seed convergence. For each subset size k , we draw n_{sub} random subsets of k seeds from the full 100-seed pool, average their predictions, and report the resulting out-of-sample L/S Sharpe ratio. Mean Sharpe stabilizes by $k \approx 50$, justifying the production choice; standard deviation across subsets shrinks as $1/\sqrt{k}$.

Two patterns emerge. First, the mean Sharpe rises from 1.04 at $k = 1$ to 1.12 at $k = 100$, so a single-chain point estimate understates the strategy’s performance by roughly 0.08 in Sharpe units. The gap closes monotonically, with most of the lift realised by $k = 30$. Second, the interquartile range across random subsets collapses from 0.10 at $k = 1$ to 0.02 at $k = 50$: once $k \geq 50$, any two random subsets of the seed pool produce nearly the same XGBoost portfolio. The production configuration ($N_s = 200$) sits well inside this plateau, which is why the 1.11 Sharpe reported in Section 5 is insensitive to which specific seeds were drawn. The practical implication is methodological rather than cosmetic: a single-chain HMM would not reliably support the main conclusions, so the seed ensemble is a requirement of the design, not a late-stage performance tuning.

Appendix J: International HMM Feature Selection

For the international robustness check (Section 5.4), the HMM feature set is selected per region using the same 4-pass procedure as the US production specification (Section 4.2): a quality screen on MCMC convergence and crisis-window alignment, a portfolio-Sharpe ranking on the surviving combinations, a definitive validation at higher seed counts, and an additional stability check across random seed partitions. The international candidate pool comprises the five locally-available stress features (DD, VOL, DISP, REL_N, BANK_REL); BANK_REL is the value-weighted bank-sector excess return over the broad regional market, sign-flipped so high values indicate stress, and substitutes for the US-specific Moody’s BAA-AAA spread. With DD required, the procedure evaluates fifteen combinations of size one to four. Both UK and Japan select $DD+VOL+REL_N$ as the regional feature set. The transplanted-from-US template ($DD+DISP+REL_N+BANK_REL$)

fails the Pass-1 quality screen for both regions (minimum effective sample size of 14 and 4 respectively), motivating per-region selection. Pass-by-pass results are reported in `results/thesis/intl_<region>_hmm_feature_selection_pass{1,2,3,4}.csv`; chosen features manifests in `intl_<region>_hmm_chosen_features.json`.

Appendix K: Fundamental Feature Definitions

Ten firm-level characteristics are computed from Compustat quarterly data (`fundq`), linked to CRSP via the CCM bridge table. All fundamentals are lagged four months from fiscal quarter end to ensure real-time availability. Cash-flow items (`oancfy`, `capxy`) are reported year-to-date and are differenced within each fiscal year to recover the per-quarter contribution before the trailing twelve-month sum is taken.

- **Book-to-market (bm):** Book equity (`ceqq`) divided by market equity.
- **Return on equity (roe):** Income before extraordinary items (`ibq`) divided by book equity.
- **Gross profitability:** Revenue minus cost of goods sold, scaled by total assets: $(\text{revtq} - \text{cogsq})/\text{atq}$.
- **Asset growth:** Year-over-year growth in total assets: $\text{atq}_t/\text{atq}_{t-4} - 1$.
- **Leverage:** Total debt divided by total assets: $(\text{dlttq} + \text{dlcq})/\text{atq}$.
- **Earnings growth:** Year-over-year growth in earnings, transformed as $\text{sign}(g) \cdot \log(1 + |g|)$.
- **Log market equity:** Natural logarithm of market capitalisation (price $\times$ shares outstanding), lagged one month.
- **Operating cash flow over assets (cfo_a):** Trailing twelve-month operating cash flow (`oancfy`, converted to quarterly and summed) scaled by total assets.
- **Free cash flow over assets (fcf_a):** Trailing twelve-month operating cash flow minus trailing twelve-month capital expenditure (`capxy`), scaled by total assets.
- **Accruals:** The accrual component of earnings, $(\text{ibq}_{\text{TMM}} - \text{oancf}_{\text{TMM}})/\text{atq}$, which captures the non-cash portion of reported earnings.

Fundamentals are winsorised at the 1st and 99th percentiles (5th and 95th for asset growth) using training-sample bounds. The cash-flow features (`cfo_a`, `fcf_a`, `accruals`) are available for roughly 83% of stock-months; the remaining seven characteristics exceed 82% coverage individually, and XGBoost handles any residual missingness natively via its default-branch logic.